

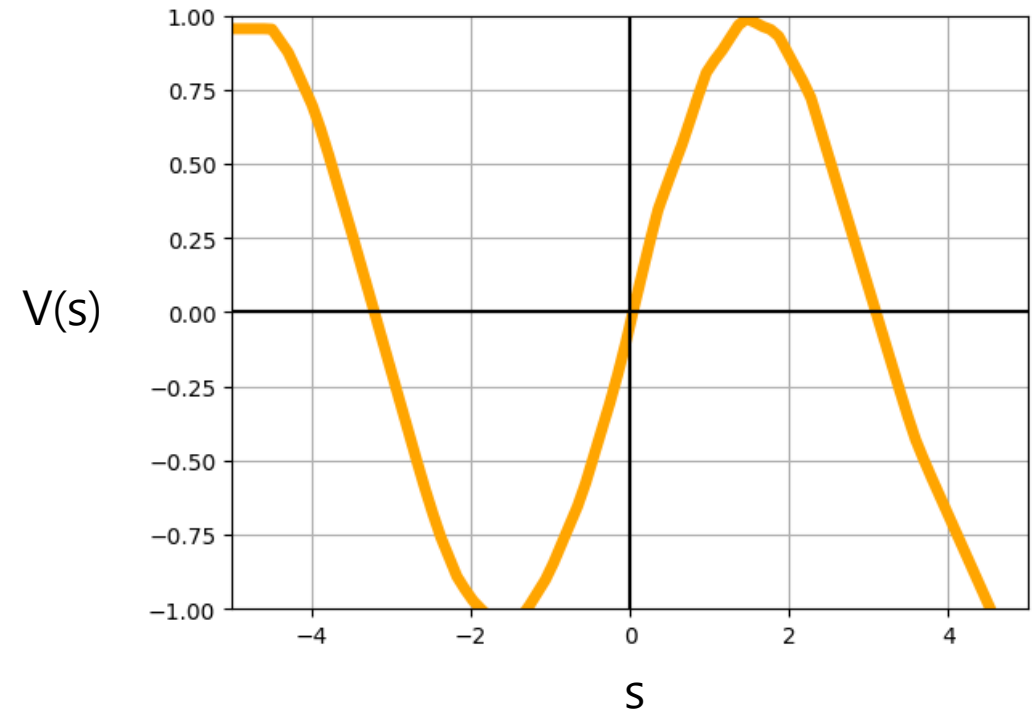
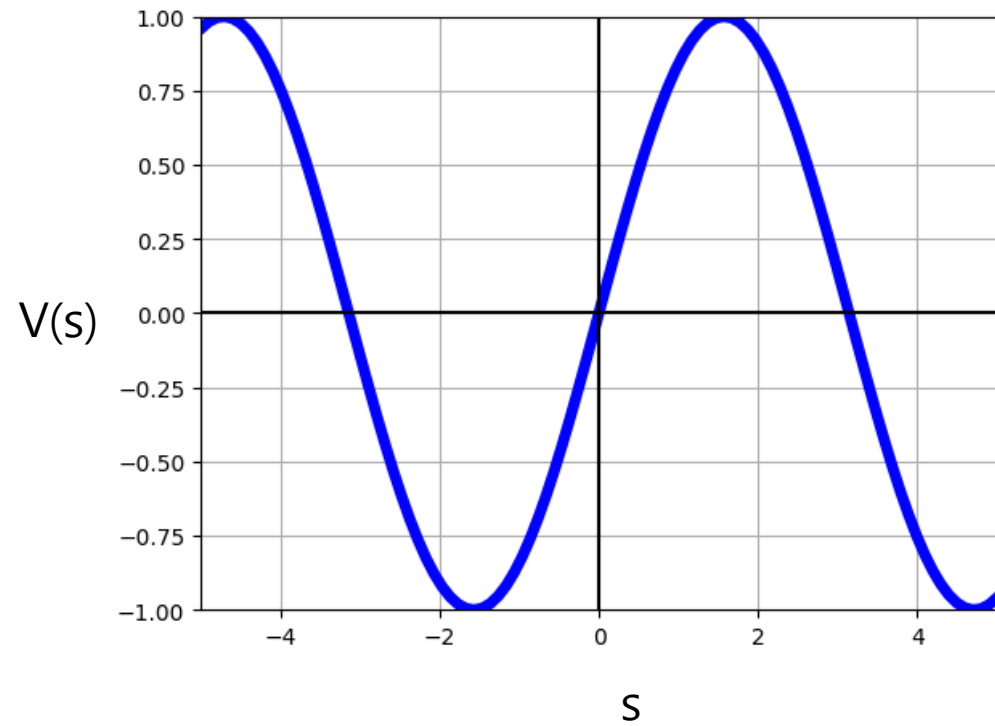
7.강 Deep Q-learning

Contents

- Neural Network
- Pytorch
- Deep SARSA
- Code Exercise

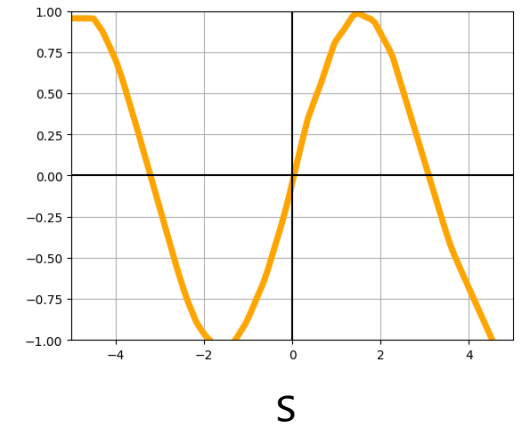
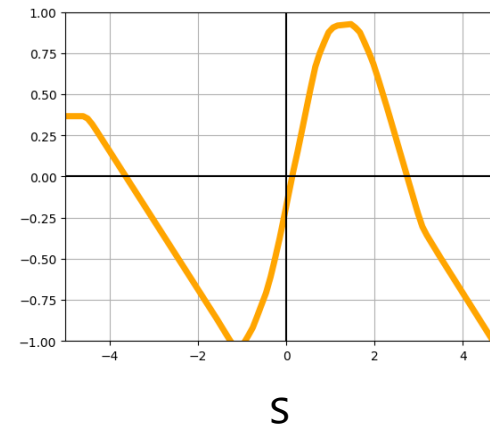
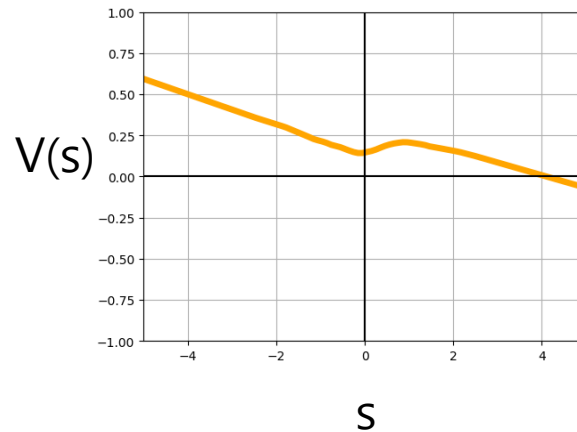
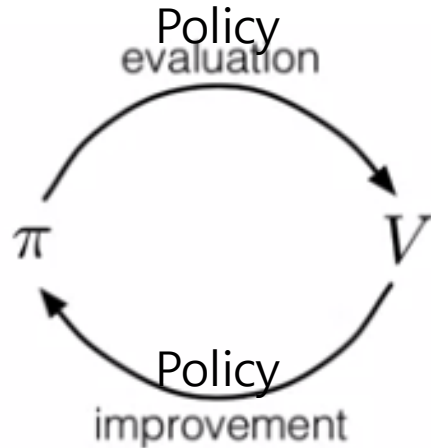
Neural Network

Function approximators



Function approximators

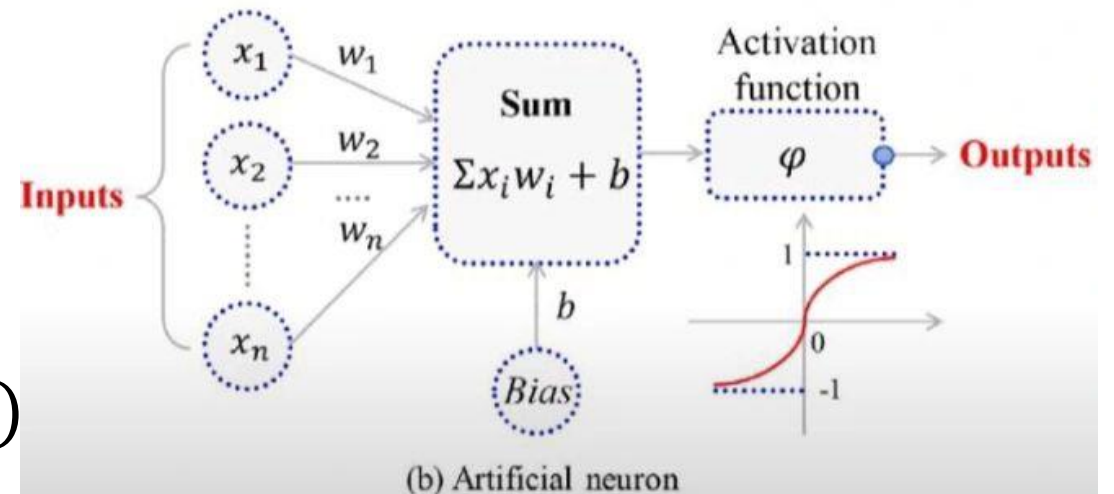
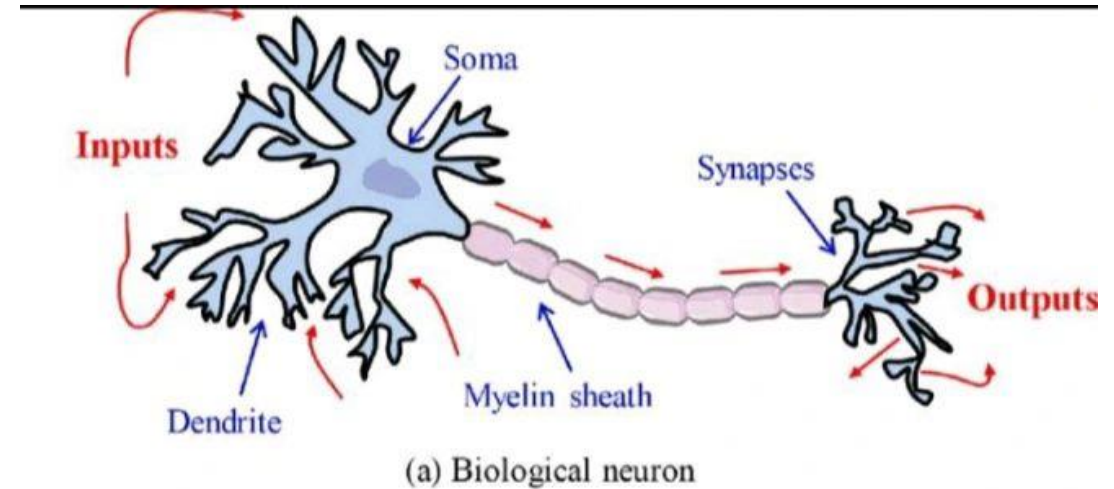
- How do we observe the value function?
 - The agent learns based on experience.
 - The functions $v^*(s)$ and $q^*(s, a)$ are not known in advance



$$f_1(s|w) \rightarrow f_2(s|w) \rightarrow \dots \rightarrow f_n(s|w) \approx v^*(s)$$

Neural Networks

- Computing system inspired by the biological neural networks that constitute our brain
- They server multiple purposes, including function approximation: $\hat{y} = f(x|w)$
- Mathematical function typically consisting of a weighted sum of inputs and a activation/transfer function
Output = $\varphi(\sum_{i=1}^n w_i x_i + b)$

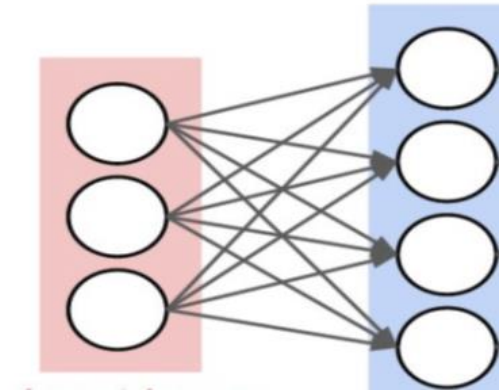


Neural networks

- Input vector $x = [x_1, x_2, x_3]$

- Connection matrix:

$$W1 = \begin{bmatrix} w_{11} & w_{12} & w_{13} & w_{14} \\ w_{21} & w_{22} & w_{23} & w_{24} \\ w_{31} & w_{32} & w_{33} & w_{34} \end{bmatrix}$$



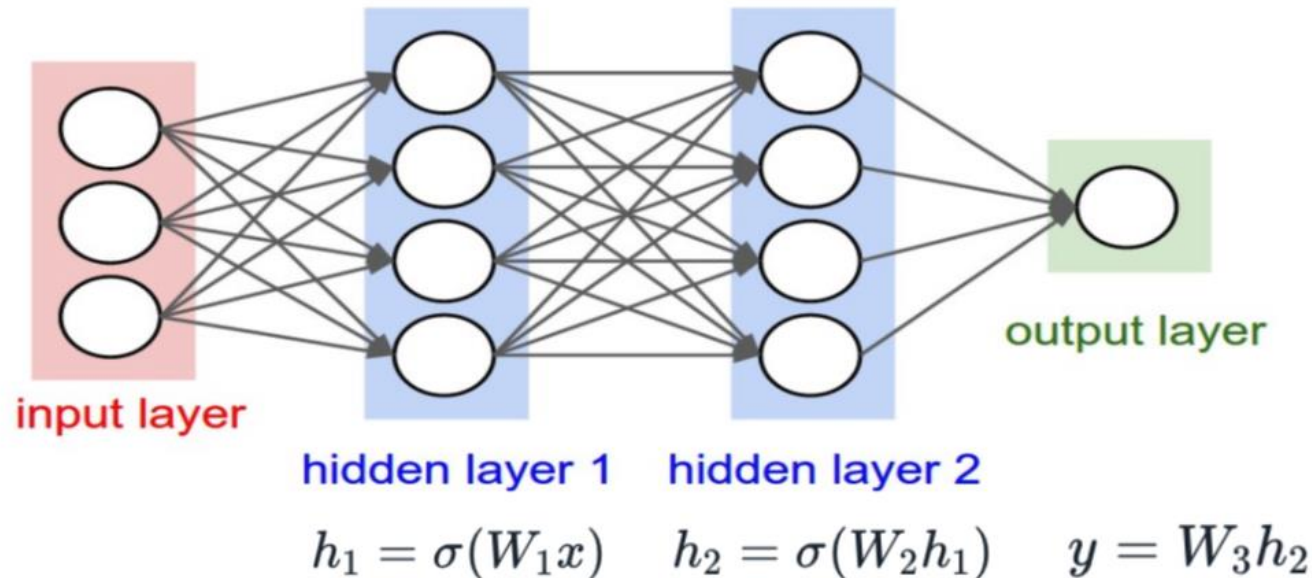
By changing its parameters $W1$, we can modify it to approximate the function we are interested in

- Output vector:

$$H = [\varphi(\sum_{i=1}^n w_i x_i + b), \dots, \varphi(\sum_{i=1}^n w_i x_i + b)]$$

Neural Networks

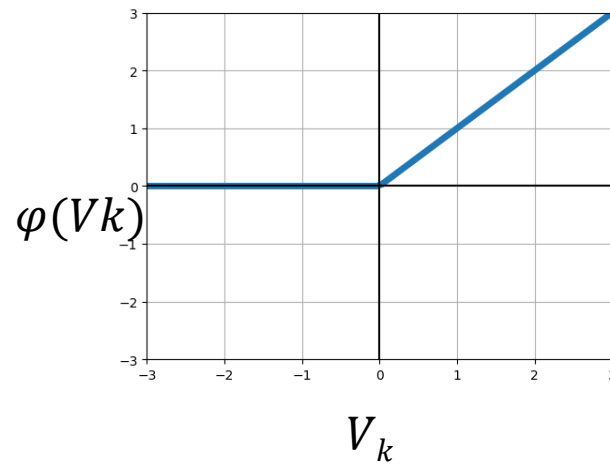
- Networks that do not have cycles are known as feedforward NN. Signals always propagate forward
- The neuron receives inputs, process & aggregate those inputs, and either inhibits or amplifies before passing the signal to the next layer



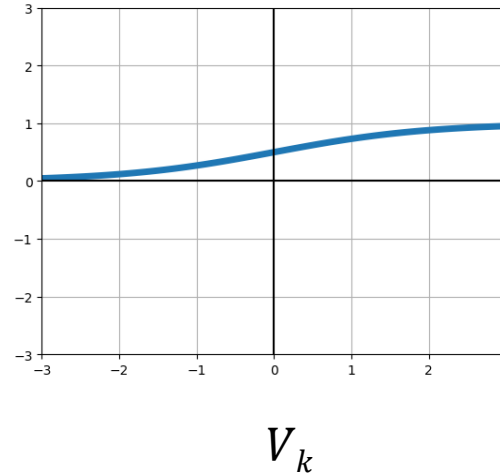
Neural networks

- Activation functions

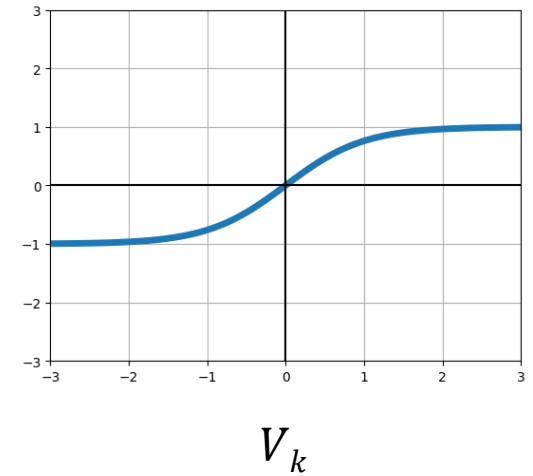
Relu()



Sigmoid()

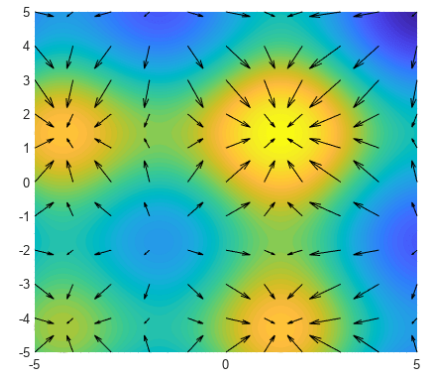
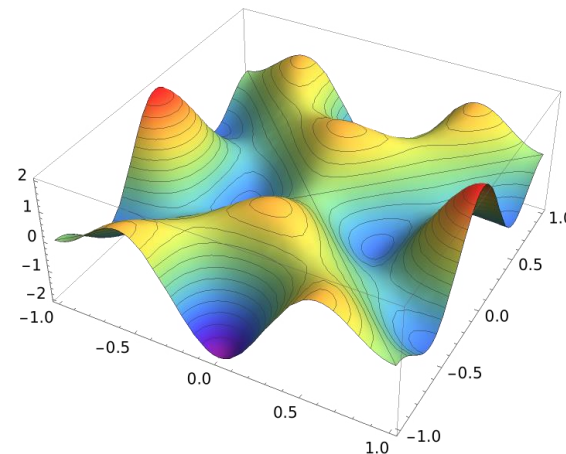
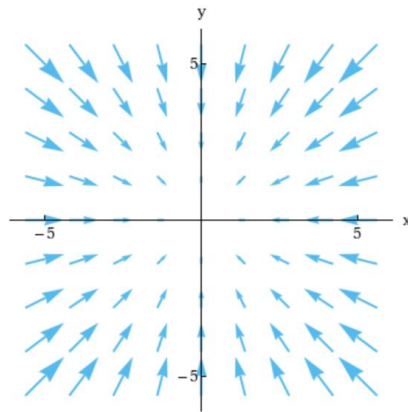
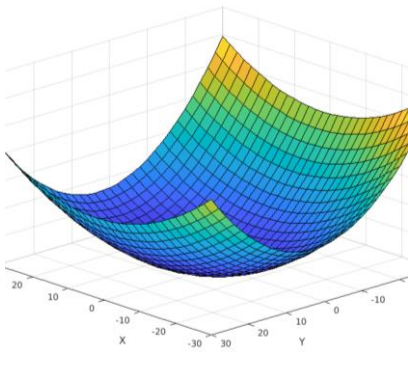


Tanh()



Gradient Descent

- Given some loss function: $L(\vec{x}, \vec{y}) = \|2\vec{x} + 2\vec{y}\|$
- Update rules for the parameters: $w_{t+1} = w_t - \alpha \nabla \hat{L}(w)$
- Gradient vector: $\nabla \hat{L}(w) = [\frac{\partial L}{\partial w_1}, \frac{\partial L}{\partial w_2}, \dots, \frac{\partial L}{\partial w_n}]$
- Computed using the backpropagation algorithm
- $\nabla \hat{L}(w)$ points to the direction of maximum growth of $\nabla \hat{L}(w)$
- α is the size of the step we take in the opposite direction to $\nabla \hat{L}(w)$



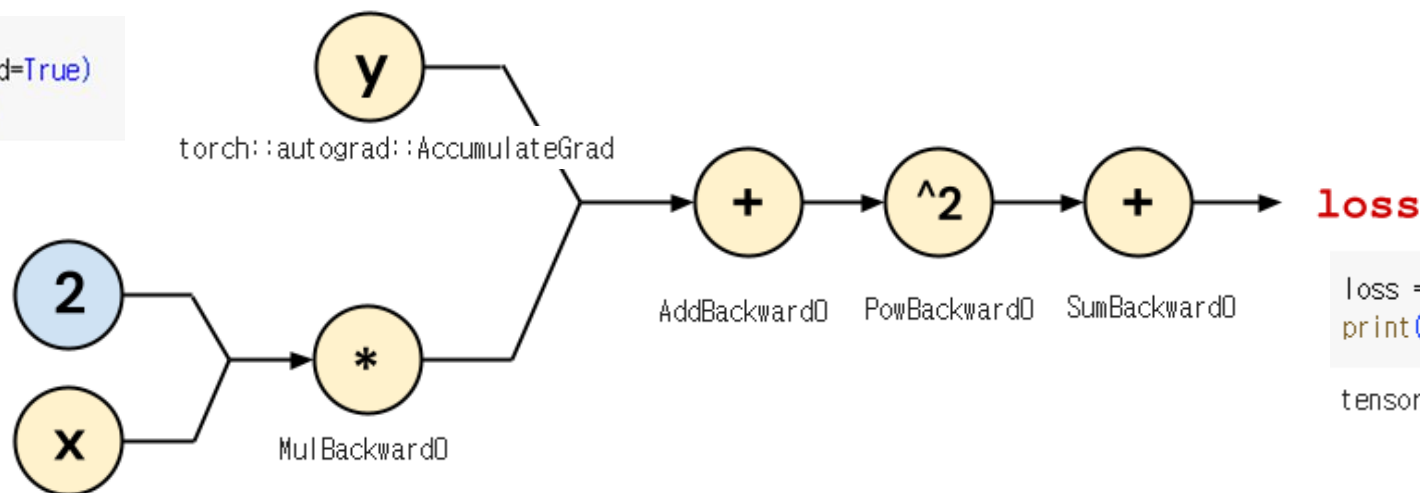
Backpropagation

- Computed using the backpropagation algorithm
- We want to evaluate partial derivative: $\frac{\partial L}{\partial \vec{x}}$ and $\frac{\partial L}{\partial \vec{y}}$

```
shape = (3, )  
x = torch.tensor([1., 2, 3], requires_grad=True)  
y = torch.ones(shape, requires_grad=True)
```

```
loss = ((2 * x + y)**2).sum()  
loss.backward()  
print(x.grad)  
print(y.grad)
```

```
tensor([24., 40., 56.])  
tensor([12., 20., 28.])
```



```
loss = ((2 * x + y)**2).sum()  
print(loss)
```

```
tensor(83., grad_fn=<SumBackward0>)
```

Cost function

- Mean squared error:

$$L(w) = \frac{1}{N} \sum_{i=0}^N [y - \hat{y}]^2$$

- For our neural network to estimate $q(s, a)$ as well as possible, we will minimize the observed squared errors

$$\hat{L}(w) = \frac{1}{N} \sum_{i=0}^N [R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}|w) - \hat{q}(S_t, A_t|w)]^2$$

- Target value:

$$R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}|w)$$

Estimated value:

$$\hat{y} = \hat{q}(S_t, A_t|w)$$

Neural network optimization

- For our neural network to estimate $q(s, a)$ as well as possible, we will minimize the observed squared errors

$$\hat{L}(\theta) = \frac{1}{N} \sum_{i=0}^N [R_i + \gamma \hat{q}(S_i', A_i' | \theta_{target}) - \hat{q}(S_i, A_i | \theta)]^2$$

- Target value: a value towards which we want to push the estimates

$$R_i + \gamma \hat{q}(S_i', A_i' | \theta_{target})$$

- Estimate of the q-value of a state-action pair

$$\hat{q}(S_i, A_i | \theta)$$

Pytorch

Pictures from Stanford's CS231n
Pictures from Berkeley CS285

Numpy & PyTorch



- Fast CPU implementations
- **CPU-only**
- **No autodiff**
- Imperative

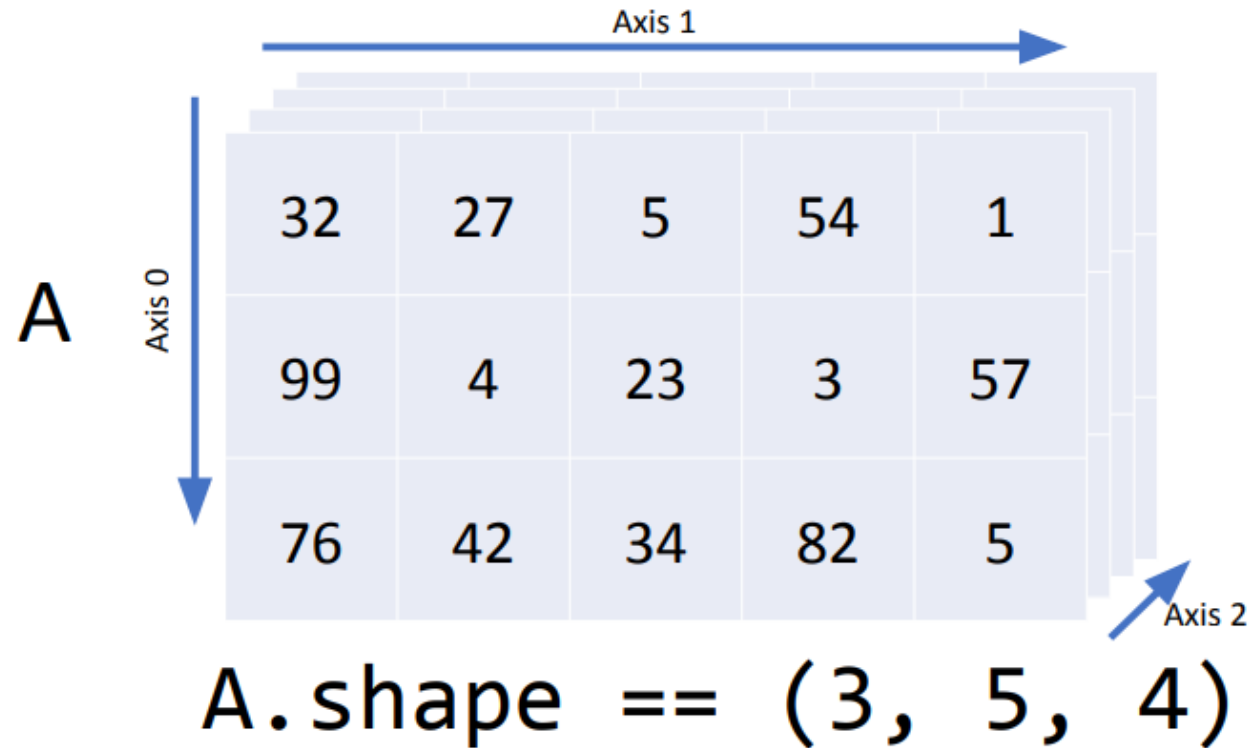


- Fast CPU implementations
- **Allows GPU**
- **Supports autodiff**
- Imperative

Other features include:

- Datasets and dataloading
- Common neural network operations
- Built-in optimizers (Adam, SGD, ...)

Multidimensional Indexing



Shape Operations



```
A = np.random.normal(size=(10, 15))

# Indexing with newaxis/None
# adds an axis with size 1
A[np.newaxis] # -> shape (1, 10, 15)

# Squeeze removes a axis with size 1
A[np.newaxis].squeeze(0) # -> shape (10, 15)

# Transpose switches out axes.
A.transpose((1, 0)) # -> shape (15, 10)

# !!! BE CAREFUL WITH RESHAPE !!!
A.reshape(15, 10) # -> shape (15, 10)
A.reshape(3, 25, -1) # -> shape (3, 25, 2)
```



```
A = torch.randn((10, 15))

# Indexing with None
# adds an axis with size 1
A[None] # -> shape (1, 10, 15)

# Squeeze removes a axis with size 1
A[None].squeeze(0) # -> shape (10, 15)

# Permute switches out axes.
A.permute((1, 0)) # -> shape (15, 10)

# !!! BE CAREFUL WITH VIEW !!!
A.view(15, 10) # -> shape (15, 10)
A.view(3, 25, -1) # -> shape (3, 25, 2)
```

Device Management

- Numpy: all arrays live on the CPU's RAM
- Torch: tensors can either live on CPU or GPU memory
 - Move to GPU with `.to("cuda")` / `.cuda()`
 - Move to CPU with `.to("cpu")` / `.cpu()`

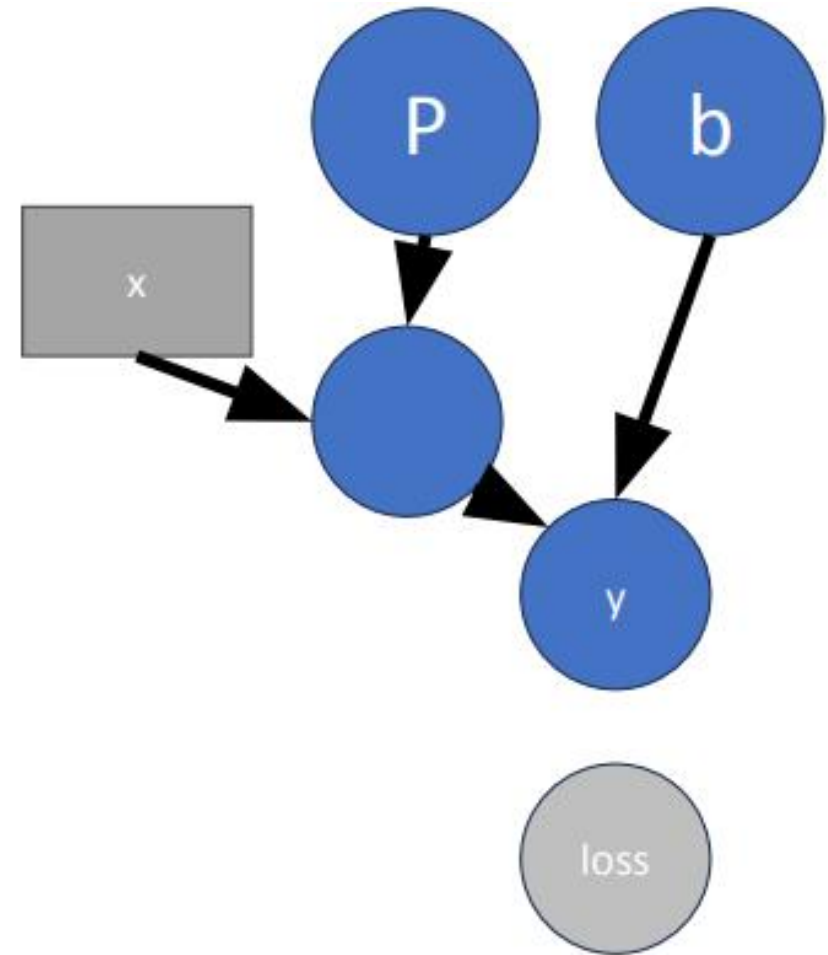
**YOU CANNOT PERFORM OPERATIONS BETWEEN
TENSORS ON DIFFERENT DEVICES!**

Computing Gradients

```
P = torch.randn((1024, 1024))
print(P.requires_grad) # -> False
P = torch.randn((1024, 1024), requires_grad=True)
b = torch.randn((1024,), requires_grad=True)
print(P.grad) # -> None

x = torch.randn((32, 1024))
y = torch.nn.relu(x @ P + b)

target = 3
loss = torch.mean((y - target) ** 2).detach()
```



Training Loop

REMEMBER THIS!

```
net = (...).to("cuda")
dataset = ...
dataloader = ..
optimizer = ...
loss_fn = ..
for epoch in range(num_epochs):
    # Training..
    net.train()
    for data, target in dataloader:
        data = torch.from_numpy(data).float().cuda()
        target = torch.from_numpy(target).float().cuda()

        prediction = net(data)
        loss = loss_fn(prediction, target)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    net.eval()
    # Do evaluation..
```

Converting Numpy / PyTorch

Numpy -> PyTorch:

```
torch.from_numpy(numpy_array).float()
```

PyTorch -> Numpy:

- (If requires_grad) Get a copy without graph with `.detach()`
- (If on GPU) Move to CPU with `.to("cpu")/.cpu()`
- Convert to numpy with `.numpy`

All together:

```
torch_tensor.detach().cpu().numpy()
```

Custom networks

```
import torch.nn as nn

class SingleLayerNetwork(nn.Module):
    def __init__(self, in_dim: int, out_dim: int, hidden_dim: int):
        super().__init__() # <- Don't forget this!
        self.net = nn.Sequential(
            nn.Module(in_dim, hidden_dim),
            nn.ReLU(),
            nn.Module(hidden_dim, out_dim),
        )

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.net(x)

batch_size = 256
my_net = SingleLayerNetwork(2, 32, 1).to("cuda")
output = my_net(torch.randn(size=(batch_size, 2)).cuda())
```

- `nn.Module` represents the building blocks of a computation graph.
 - For example, in typical pytorch code, each convolution block is its own module, each fully connected block is a module, and the whole network itself is also a module.
- Modules can contain modules within them. All the classes inside of ``torch.nn`` are instances ``nn.Modules``.

Custom networks

```
import torch.nn as nn

class SingleLayerNetwork(nn.Module):
    def __init__(self, in_dim: int, out_dim: int, hidden_dim: int):
        super().__init__() # <- Don't forget this!
        self.net = nn.Sequential(
            nn.Module(in_dim, hidden_dim),
            nn.ReLU(),
            nn.Module(hidden_dim, out_dim),
        )

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.net(x)

batch_size = 256
my_net = SingleLayerNetwork(2, 32, 1).to("cuda")
output = my_net(torch.randn(size=(batch_size, 2)).cuda())
```

- Prefer `net()` over `net.forward()`
- Everything (network and its inputs) on the same device!!!

Torch Best Practices

- When in doubt, **assert** is your friend

```
assert x.shape == (B, N), \
    f"Expected shape ({B}, {N}) but got {x.shape}"
```

- Be extra careful with **.reshape/.view**
 - If you use it, assert before and after
 - Only use it to collapse/expand a single dim
 - In Torch, prefer **.flatten()/.permute()/.unflatten()**
- If you do some complicated operation, test it!
 - Compare to a pure Python implementation

Torch Best Practices

- Don't mix numpy and Torch code
 - Understand the boundaries between the two
 - Make sure to cast 64-bit numpy arrays to 32 bits
 - `torch.Tensor` only in `nn.Module`!
- Training loop will always look the same
 - Load batch, compute loss
 - `.zero_grad()`, `.backward()`, `.step()`

Neural network optimization

- Mean squared error:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N [y_i - \hat{y}_i]^2$$

- We want to minimize the square of the errors of the neural network estimates

Neural network optimization

- We calculate the gradient vector of the cost function with respect to the θ parameters:

$$\nabla L(\theta) = \left[\frac{\partial L}{\partial \theta_1}, \frac{\partial L}{\partial \theta_2}, \dots, \frac{\partial L}{\partial \theta_n} \right]$$

- With the gradient vector, we will make a SGD step:

$$\theta \leftarrow \theta - \alpha \nabla \hat{L}(\theta)$$

Neural Net. Architecture for V, Q

- St vector input \rightarrow NN \rightarrow V scalar output
- St, At vector input \rightarrow NN \rightarrow Q scalar output
 - Continuous case
- St vector input \rightarrow Q vector output
 - Discrete action space only
 - Output size is $|A|$

Neural Net. Architecture for policy, π

- S_t input \rightarrow NN \rightarrow vector output
 - Discrete action space case
 - Output size is $|A|$
 - SOFTMAX turns the output into prob. (sum of prob. is 1)
- S_t input \rightarrow NN $\rightarrow \mu_{\theta}(S_t), \delta_{\theta}(S_t)$ output
 - Continuous action space case
 - Represented with Gaussian Distribution

Deep SARSA

Neural network optimization

$$L(\theta) = \frac{1}{|K|} \sum_{i=0}^{|K|} [R_i + \gamma \hat{q}(S'_i, A'_i | \theta_{target}) - \hat{q}(S_i, A_i | \theta)]^2$$

- Target is the value towards which we want to push the estimates.

$$R_i + \gamma \hat{q}(S'_i, A'_i | \theta_{target})$$

- Estimate is the estimate of the q-value of a state-action pair
 $\hat{q}(S_i, A_i | \theta)$

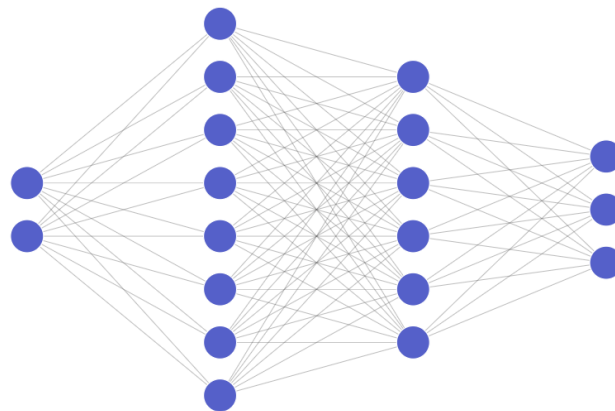
Target network

Bootstrapping



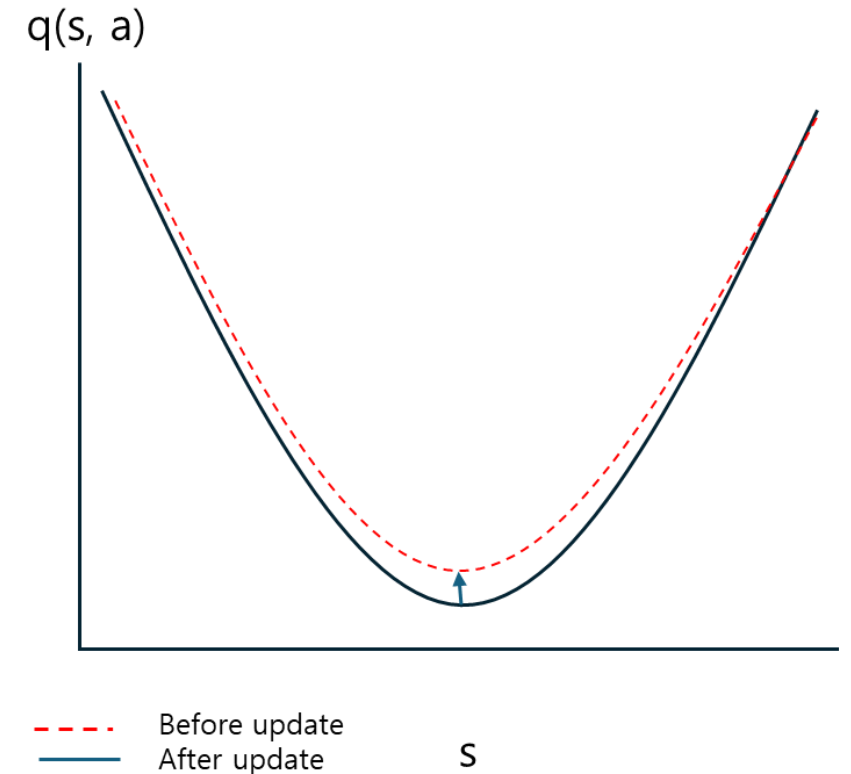
Function
approximator

$$y_i = R_i + \gamma \hat{q}(S'_i, A'_i | \theta_{target})$$



Target network

- When a value is changed, nearby values will also be affected.
- By modifying a $\hat{q}(S_i, A_i | \theta)$ estimate we also modify its $\hat{q}(S_i', A_i' | \theta_{target})$ target
- For the learning process to be stable, the target must also be stable
- power of neural networks



Target network

- We make a copy of the neural network to calculate the targets.

$$\theta_{targ} \leftarrow \theta$$

- This neural network does not change with SGD. Its θ parameters remain the same
- The estimated value of S_i', A_i' is calculated with the target network:

$$L(\theta) = \frac{1}{N} \sum_{i=0}^N [R_i + \gamma \hat{q}(S_i', A_i' | \theta_{targ}) - \hat{q}(S_i, A_i | \theta)]^2$$

Neural network optimization

Algorithm 1 Deep SARSA

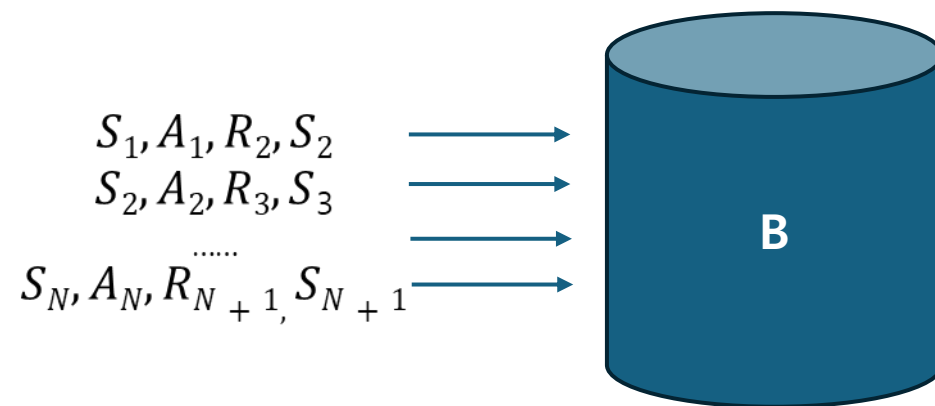
- 1: **Input:** α learning rate, ϵ random action probability,
- 2: γ discount factor,
- 3: Initialize q-value parameters θ and target parameters $\theta_{targ} \leftarrow \theta$
- 4: $\pi \leftarrow \epsilon$ -greedy policy w.r.t $\hat{q}(s, a|\theta)$
- 5: Initialize replay buffer B
- 6: **for** episode $\in 1..N$ **do**
- 7: Restart environment and observe the initial state S_0
- 8: **for** $t \in 0..T - 1$ **do**
- 9: Select action $A_t \sim \pi(S_t)$
- 10: Execute action A_t and observe S_{t+1}, R_{t+1}
- 11: Insert transition $(S_t, A_t, R_{t+1}, S_{t+1})$ into the buffer B
- 12: $K = (S, A, R, S') \sim B$
- 13: Select actions $A' \sim \pi(S')$
- 14: Compute loss function over the batch of experiences:

$$L(\theta) = \frac{1}{|K|} \sum_{i=1}^{|K|} [R_i + \gamma \hat{q}(S'_i, A'_i|\theta_{targ}) - \hat{q}(S_i, A_i|\theta)]^2 \quad (1)$$

- 15: **end for**
 - 16: Every k episodes synchronize $\theta_{targ} \leftarrow \theta$
 - 17: **end for**
 - 18: **Output:** Near optimal policy π and q-value approximations $\hat{q}(s, a|\theta)$
-

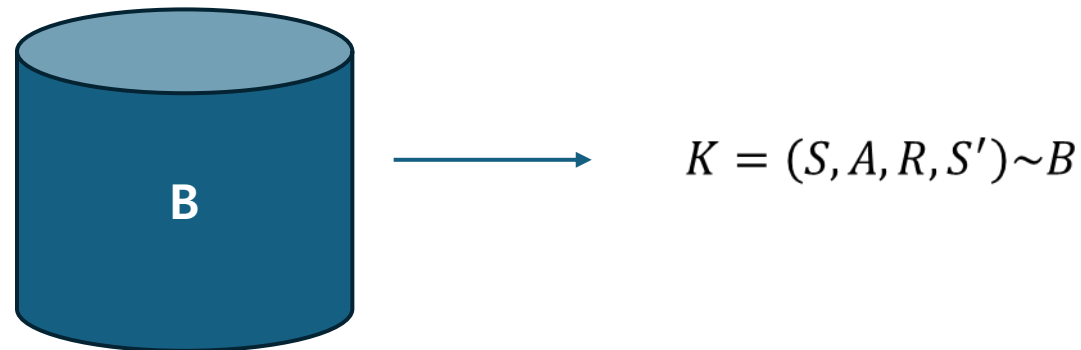
Experience Replay

- Memory that stores the state transition that the agent experiences
- The memory has a limited size and when it fills up, it replaces old transitions with new ones



Experience Replay

- To update the neural network, we randomly chose a batch of transitions from the memory



- The batch of transitions obtained from the memory is used to calculate the cost function and update the θ parameters
- $$L(\theta) = \frac{1}{|K|} \sum_{i=0}^{|K|} [R_i + \gamma \hat{q}(S'_i, A'_i | \theta_{target}) - \hat{q}(S_i, A_i | \theta)]^2$$

Neural network optimization

Algorithm 1 Deep SARSA

- 1: **Input:** α learning rate, ϵ random action probability,
 - 2: γ discount factor,
 - 3: Initialize q-value parameters θ and target parameters $\theta_{target} \leftarrow \theta$
 - 4: $\pi \leftarrow \epsilon$ -greedy policy w.r.t $\hat{q}(s, a|\theta)$
 - 5: Initialize replay buffer B
 - 6: **for** episode $\in 1..N$ **do**
 - 7: Restart environment and observe the initial state S_0
 - 8: **for** $t \in 0..T - 1$ **do**
 - 9: Select action $A_t \sim \pi(S_t)$
 - 10: Execute action A_t and observe S_{t+1}, R_{t+1}
 - 11: Insert transition $(S_t, A_t, R_{t+1}, S_{t+1})$ into the buffer B
 - 12: $K = (S, A, R, S') \sim B$
 - 13: Select actions $A' \sim \pi(S')$
 - 14: Compute loss function over the batch of experiences:

$$L(\theta) = \frac{1}{|K|} \sum_{i=1}^{|K|} [R_i + \gamma \hat{q}(S'_i, A'_i|\theta_{target}) - \hat{q}(S_i, A_i|\theta)]^2 \quad (1)$$
 - 15: **end for**
 - 16: Every k episodes synchronize $\theta_{target} \leftarrow \theta$
 - 17: **end for**
 - 18: **Output:** Near optimal policy π and q-value approximations $\hat{q}(s, a|\theta)$
-

Code Ex.

- MountainCar: Reach the goal from the bottom of the valley

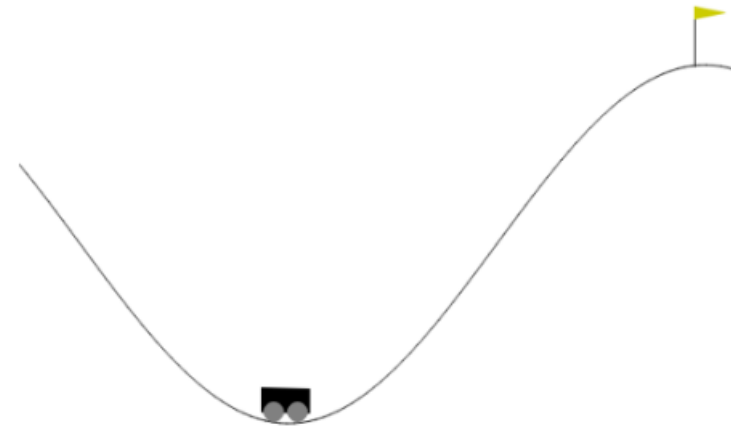
The state

The observation space consists of the car position $\in [-1.2, 0.6]$ and car velocity $\in [-0.07, 0.07]$

The actions available

The actions available three:

- 0 Accelerate to the left.
- 1 Don't accelerate.
- 2 Accelerate to the right.



Code Ex.

- `deep_sarsa_colab.ipynb`