

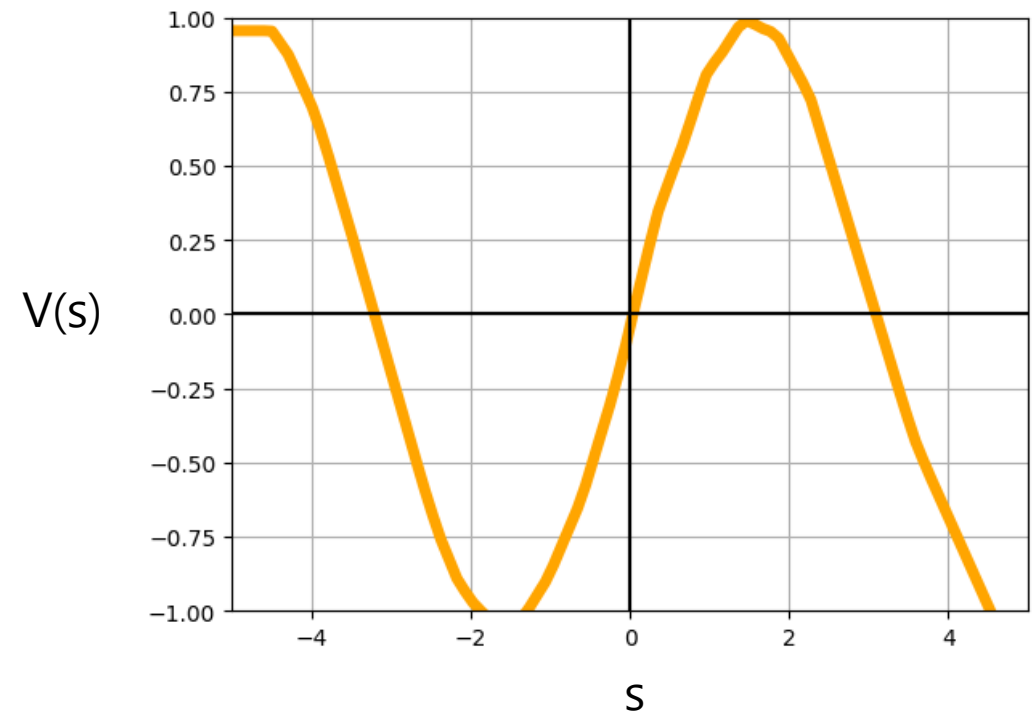
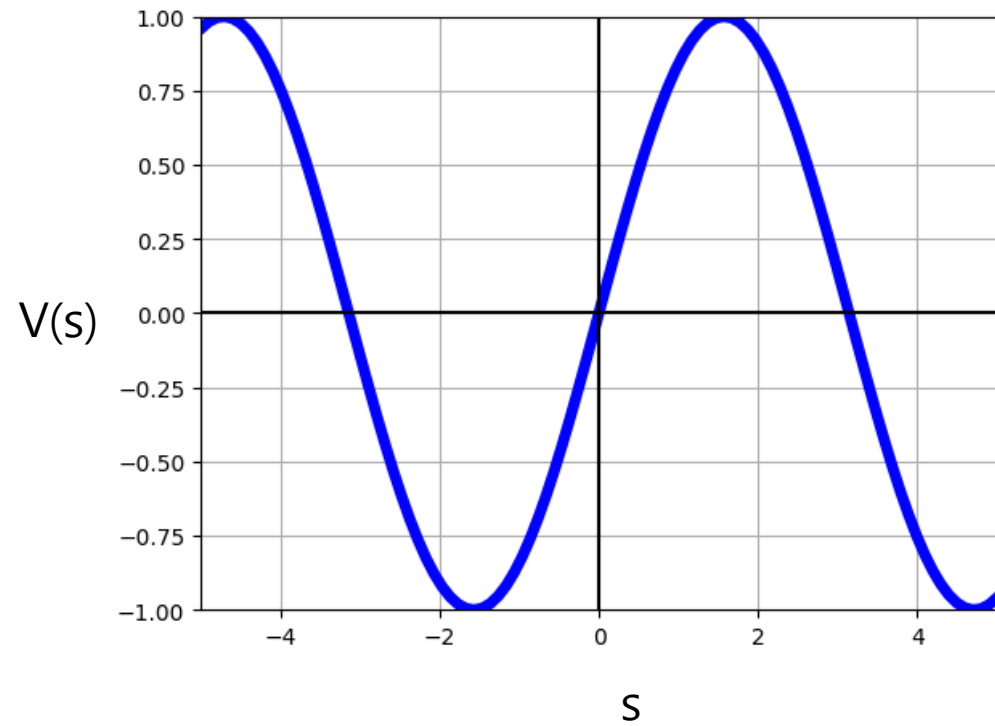
7.강 Deep Q-learning

Contents

- Neural Network
- Pytorch
- Deep Q-Learning
- Code Exercise

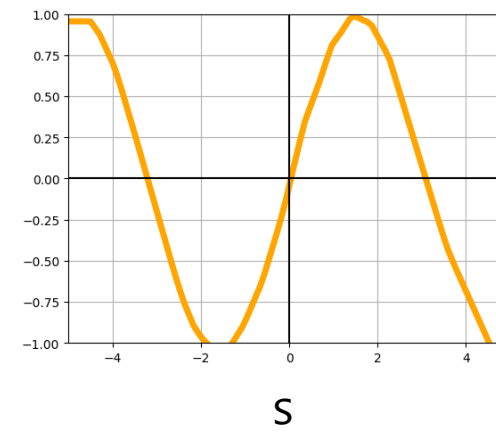
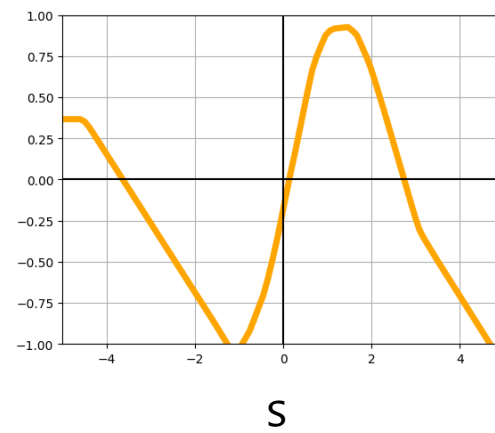
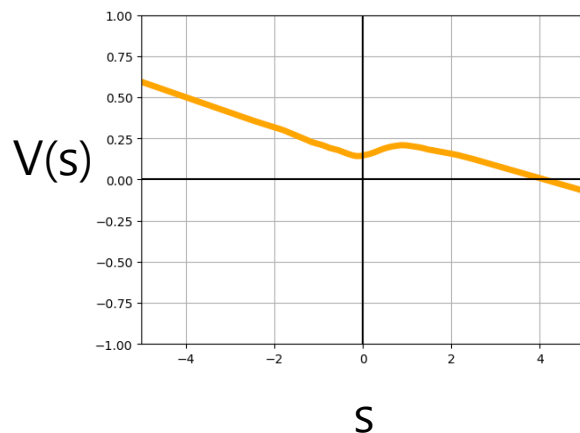
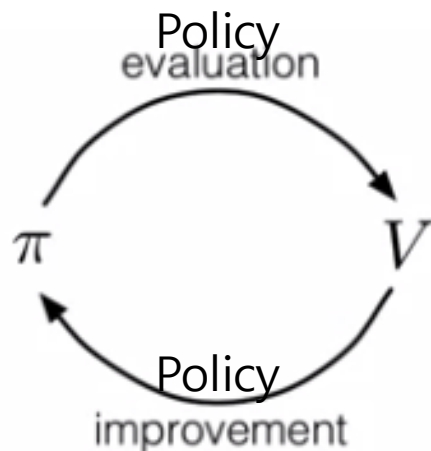
Neural Network

Function approximators



Function approximators

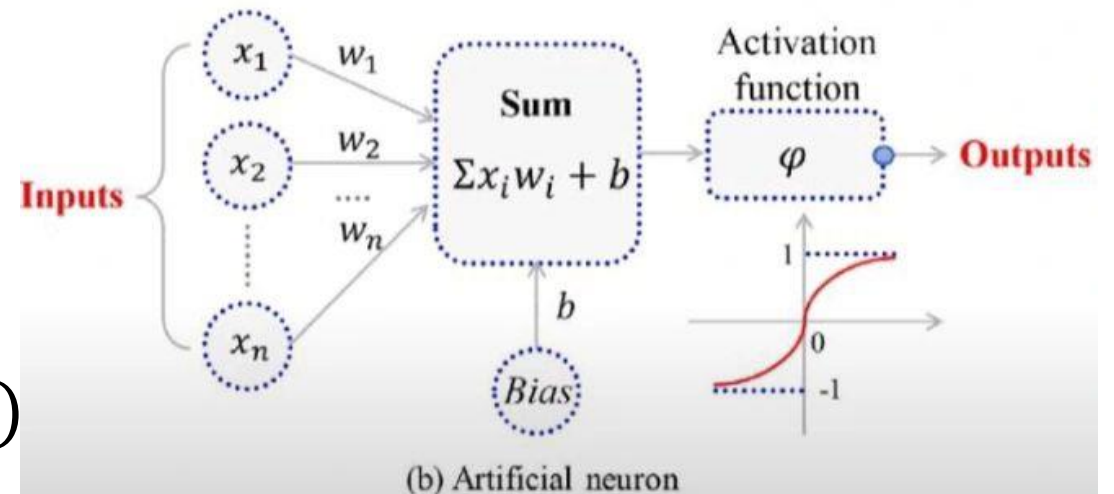
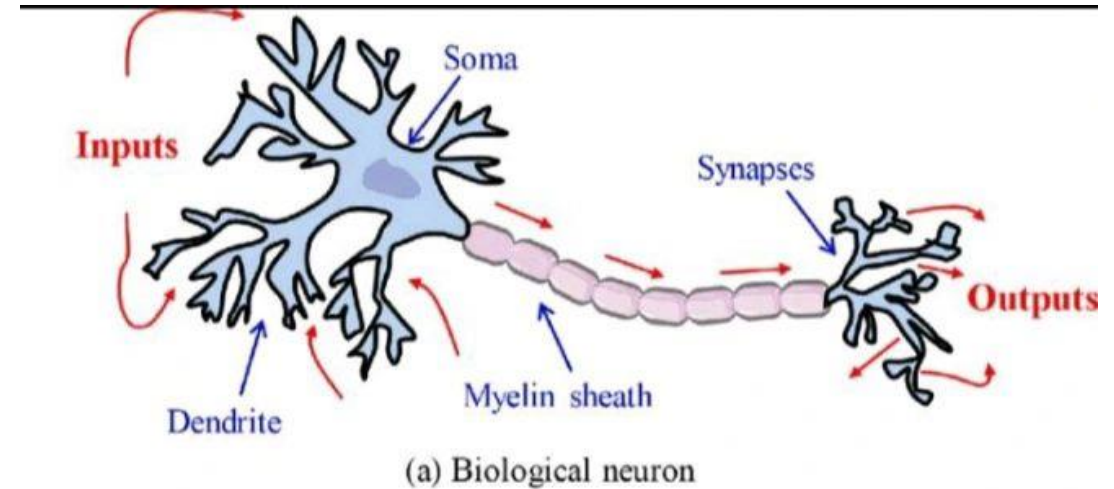
- How do we observe the value function?
 - The agent learns based on experience.
 - The functions $v^*(s)$ and $q^*(s, a)$ are not known in advance



$$f_1(s|w) \rightarrow f_2(s|w) \rightarrow \dots \rightarrow f_n(s|w) \approx v^*(s)$$

Neural Networks

- Computing system inspired by the biological neural networks that constitute our brain
- They server multiple purposes, including function approximation: $\hat{y} = f(x|w)$
- Mathematical function typically consisting of a weighted sum of inputs and a activation/transfer function
Output = $\varphi(\sum_{i=1}^n w_i x_i + b)$

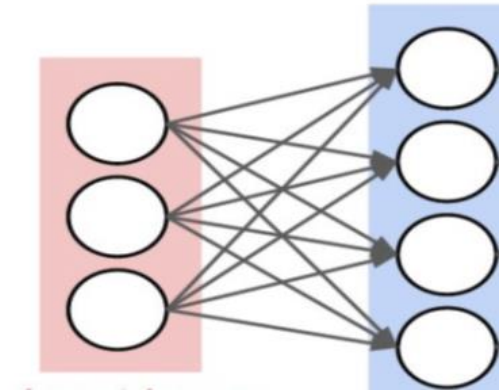


Neural networks

- Input vector $x = [x_1, x_2, x_3]$

- Connection matrix:

$$W1 = \begin{bmatrix} w_{11} & w_{12} & w_{13} & w_{14} \\ w_{21} & w_{22} & w_{23} & w_{24} \\ w_{31} & w_{32} & w_{33} & w_{34} \end{bmatrix}$$



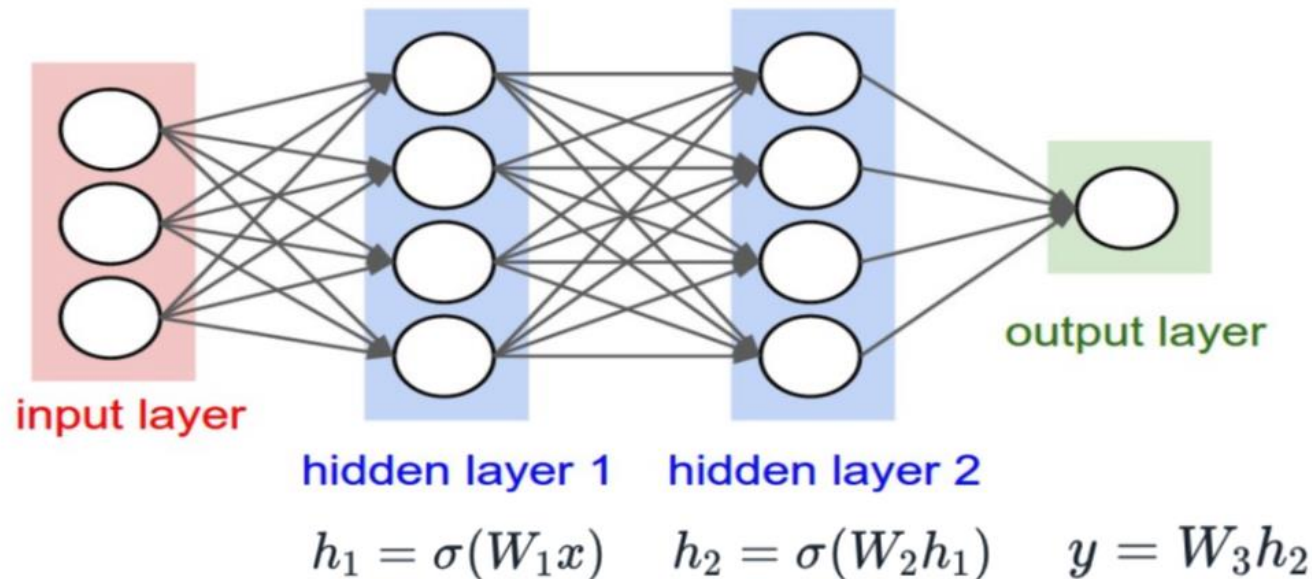
By changing its parameters $W1$, we can modify it to approximate the function we are interested in

- Output vector:

$$H = [\varphi(\sum_{i=1}^n w_i x_i + b), \dots, \varphi(\sum_{i=1}^n w_i x_i + b)]$$

Neural Networks

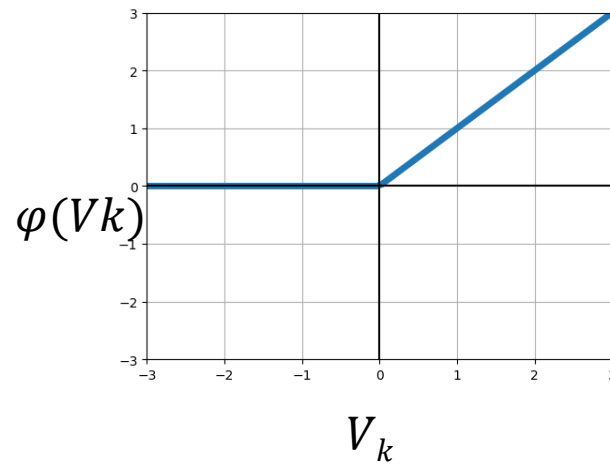
- Networks that do not have cycles are known as feedforward NN. Signals always propagate forward
- The neuron receives inputs, process & aggregate those inputs, and either inhibits or amplifies before passing the signal to the next layer



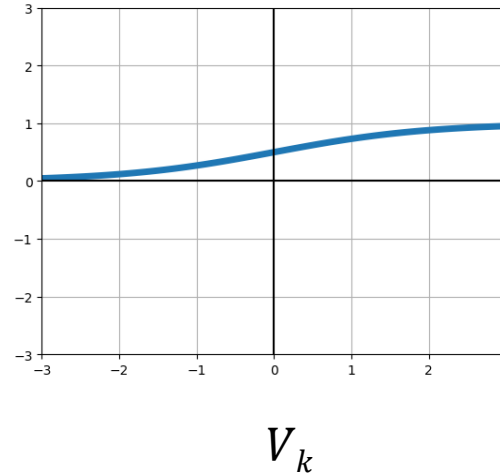
Neural networks

- Activation functions

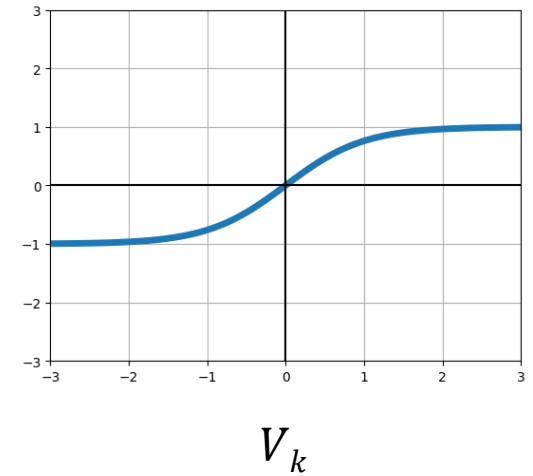
Relu()



Sigmoid()

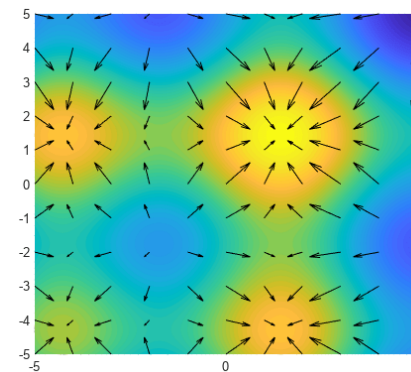
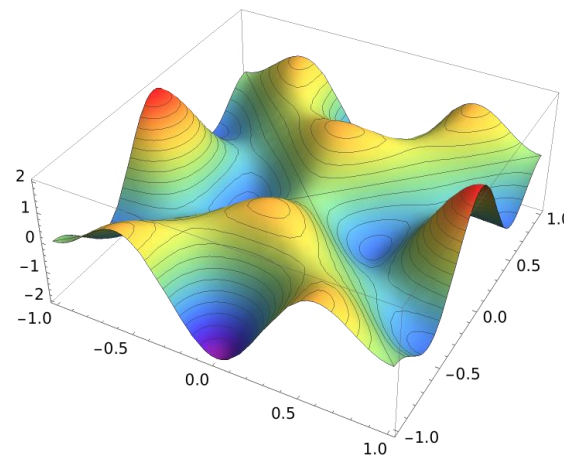
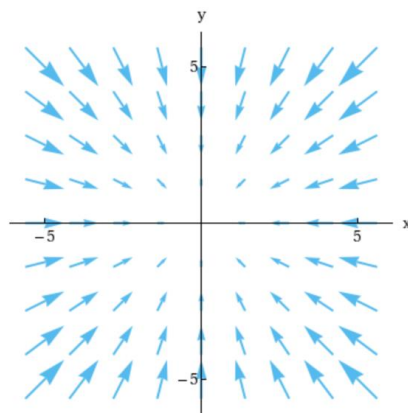
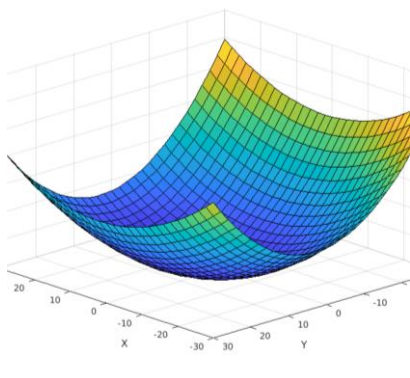


Tanh()



Gradient Descent

- Given some loss function: $L(\vec{x}, \vec{y}) = \|2\vec{x} + 2\vec{y}\|$
- Update rules for the parameters: $w_{t+1} = w_t - \alpha \nabla \hat{L}(w)$
- Gradient vector: $\nabla \hat{L}(w) = [\frac{\partial L}{\partial w_1}, \frac{\partial L}{\partial w_2}, \dots, \frac{\partial L}{\partial w_n}]$
- Computed using the backpropagation algorithm
- $\nabla \hat{L}(w)$ points to the direction of maximum growth of $\nabla \hat{L}(w)$
- α is the size of the step we take in the opposite direction to $\nabla \hat{L}(w)$



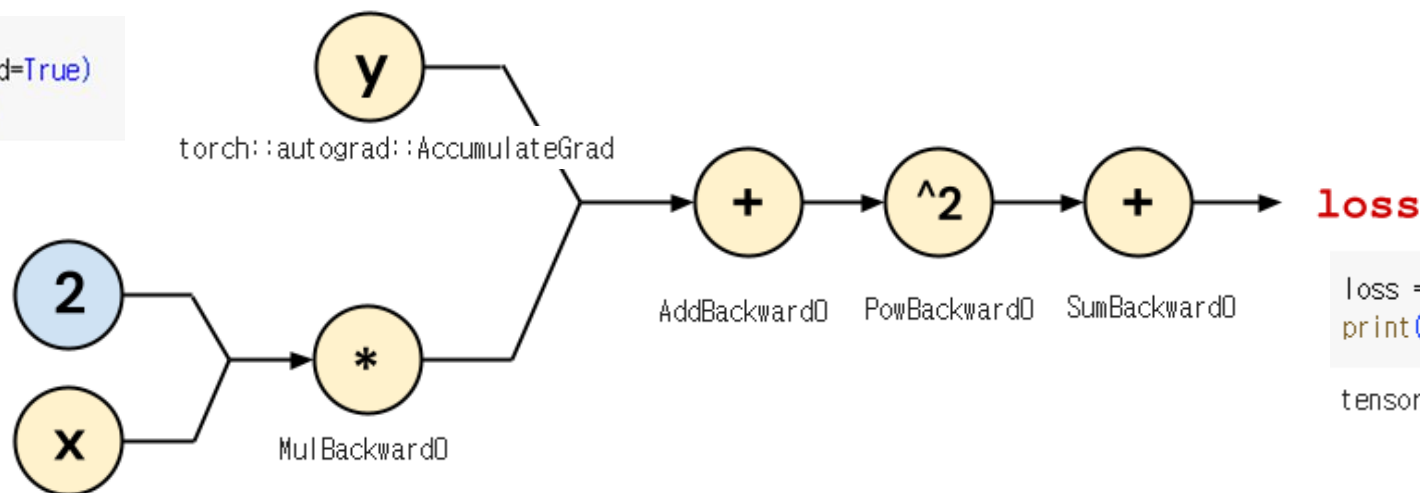
Backpropagation

- Computed using the backpropagation algorithm
- We want to evaluate partial derivative: $\frac{\partial L}{\partial \vec{x}}$ and $\frac{\partial L}{\partial \vec{y}}$

```
shape = (3, )  
x = torch.tensor([1., 2, 3], requires_grad=True)  
y = torch.ones(shape, requires_grad=True)
```

```
loss = ((2 * x + y)**2).sum()  
loss.backward()  
print(x.grad)  
print(y.grad)
```

```
tensor([24., 40., 56.])  
tensor([12., 20., 28.])
```



```
loss = ((2 * x + y)**2).sum()  
print(loss)
```

```
tensor(83., grad_fn=<SumBackward0>)
```

Cost function

- Mean squared error:

$$L(w) = \frac{1}{N} \sum_{i=0}^N [y - \hat{y}]^2$$

- For our neural network to estimate $q(s, a)$ as well as possible, we will minimize the observed squared errors

$$\hat{L}(w) = \frac{1}{N} \sum_{i=0}^N [R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}|w) - \hat{q}(S_t, A_t|w)]^2$$

- Target value:

$$R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}|w)$$

Estimated value:

$$\hat{y} = \hat{q}(S_t, A_t|w)$$

Neural network optimization

- For our neural network to estimate $q(s, a)$ as well as possible, we will minimize the observed squared errors

$$\hat{L}(\theta) = \frac{1}{N} \sum_{i=0}^N [R_i + \gamma \hat{q}(S_i', A_i' | \theta_{target}) - \hat{q}(S_i, A_i | \theta)]^2$$

- Target value: a value towards which we want to push the estimates

$$R_i + \gamma \hat{q}(S_i', A_i' | \theta_{target})$$

- Estimate of the q-value of a state-action pair

$$\hat{q}(S_i, A_i | \theta)$$

Pytorch

Pictures from Stanford's CS231n
Pictures from Berkeley CS285

Numpy & PyTorch



- Fast CPU implementations
- **CPU-only**
- **No autodiff**
- Imperative

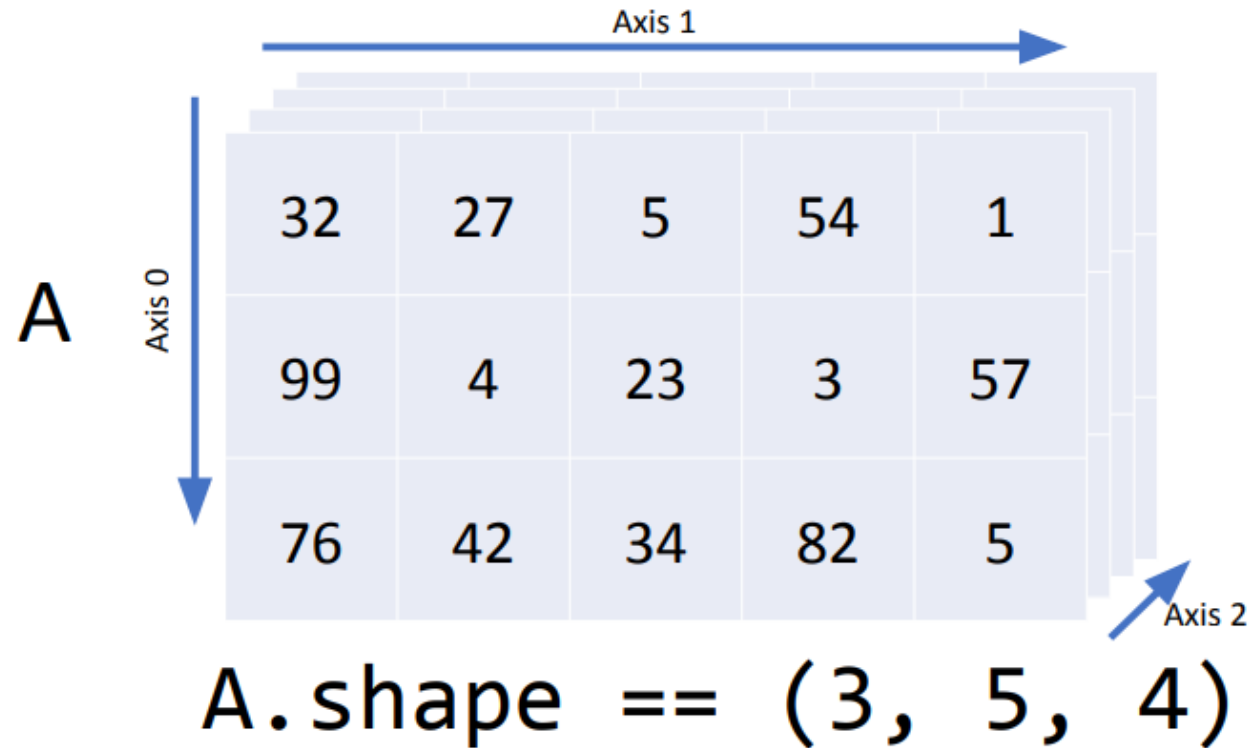


- Fast CPU implementations
- **Allows GPU**
- **Supports autodiff**
- Imperative

Other features include:

- Datasets and dataloading
- Common neural network operations
- Built-in optimizers (Adam, SGD, ...)

Multidimensional Indexing



Shape Operations



```
A = np.random.normal(size=(10, 15))

# Indexing with newaxis/None
# adds an axis with size 1
A[np.newaxis] # -> shape (1, 10, 15)

# Squeeze removes a axis with size 1
A[np.newaxis].squeeze(0) # -> shape (10, 15)

# Transpose switches out axes.
A.transpose((1, 0)) # -> shape (15, 10)

# !!! BE CAREFUL WITH RESHAPE !!!
A.reshape(15, 10) # -> shape (15, 10)
A.reshape(3, 25, -1) # -> shape (3, 25, 2)
```



```
A = torch.randn((10, 15))

# Indexing with None
# adds an axis with size 1
A[None] # -> shape (1, 10, 15)

# Squeeze removes a axis with size 1
A[None].squeeze(0) # -> shape (10, 15)

# Permute switches out axes.
A.permute((1, 0)) # -> shape (15, 10)

# !!! BE CAREFUL WITH VIEW !!!
A.view(15, 10) # -> shape (15, 10)
A.view(3, 25, -1) # -> shape (3, 25, 2)
```

Device Management

- Numpy: all arrays live on the CPU's RAM
- Torch: tensors can either live on CPU or GPU memory
 - Move to GPU with `.to("cuda")` / `.cuda()`
 - Move to CPU with `.to("cpu")` / `.cpu()`

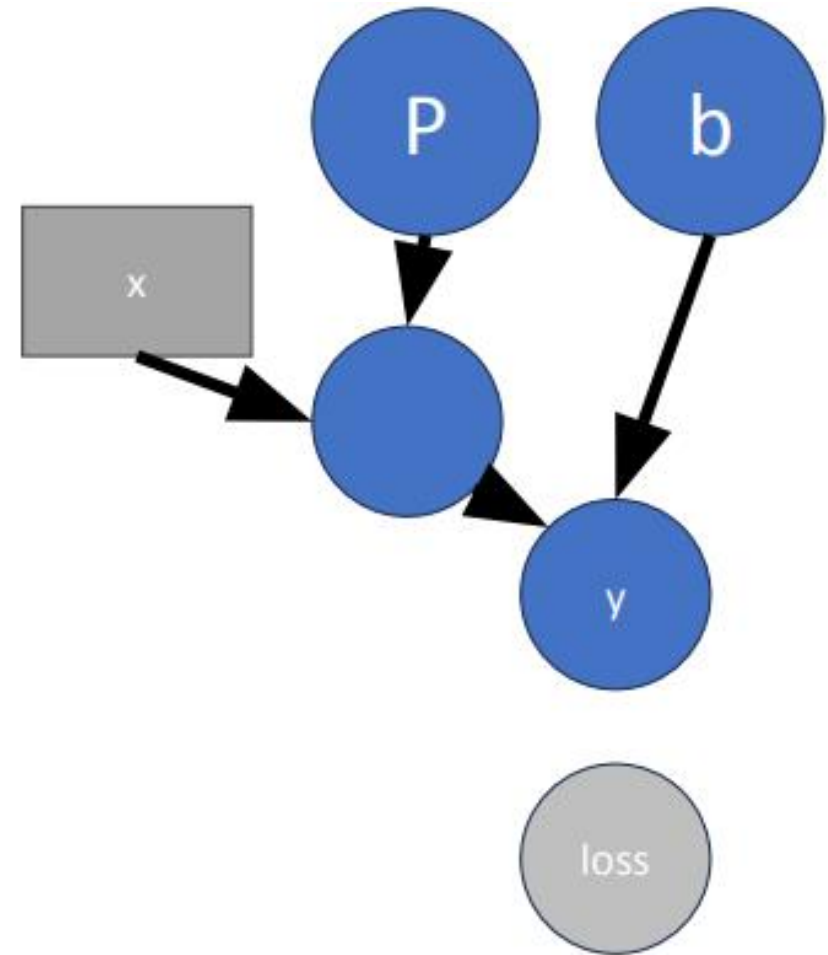
**YOU CANNOT PERFORM OPERATIONS BETWEEN
TENSORS ON DIFFERENT DEVICES!**

Computing Gradients

```
P = torch.randn((1024, 1024))
print(P.requires_grad) # -> False
P = torch.randn((1024, 1024), requires_grad=True)
b = torch.randn((1024,), requires_grad=True)
print(P.grad) # -> None

x = torch.randn((32, 1024))
y = torch.nn.relu(x @ P + b)

target = 3
loss = torch.mean((y - target) ** 2).detach()
```



Training Loop

REMEMBER THIS!

```
net = (...).to("cuda")
dataset = ...
dataloader = ..
optimizer = ...
loss_fn = ..
for epoch in range(num_epochs):
    # Training..
    net.train()
    for data, target in dataloader:
        data = torch.from_numpy(data).float().cuda()
        target = torch.from_numpy(target).float().cuda()

        prediction = net(data)
        loss = loss_fn(prediction, target)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    net.eval()
    # Do evaluation..
```

Converting Numpy / PyTorch

Numpy -> PyTorch:

```
torch.from_numpy(numpy_array).float()
```

PyTorch -> Numpy:

- (If requires_grad) Get a copy without graph with `.detach()`
- (If on GPU) Move to CPU with `.to("cpu")/.cpu()`
- Convert to numpy with `.numpy`

All together:

```
torch_tensor.detach().cpu().numpy()
```

Custom networks

```
import torch.nn as nn

class SingleLayerNetwork(nn.Module):
    def __init__(self, in_dim: int, out_dim: int, hidden_dim: int):
        super().__init__() # <- Don't forget this!
        self.net = nn.Sequential(
            nn.Module(in_dim, hidden_dim),
            nn.ReLU(),
            nn.Module(hidden_dim, out_dim),
        )

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.net(x)

batch_size = 256
my_net = SingleLayerNetwork(2, 32, 1).to("cuda")
output = my_net(torch.randn(size=(batch_size, 2)).cuda())
```

- `nn.Module` represents the building blocks of a computation graph.
 - For example, in typical pytorch code, each convolution block is its own module, each fully connected block is a module, and the whole network itself is also a module.
- Modules can contain modules within them. All the classes inside of ``torch.nn`` are instances ``nn.Modules``.

Custom networks

```
import torch.nn as nn

class SingleLayerNetwork(nn.Module):
    def __init__(self, in_dim: int, out_dim: int, hidden_dim: int):
        super().__init__() # <- Don't forget this!
        self.net = nn.Sequential(
            nn.Module(in_dim, hidden_dim),
            nn.ReLU(),
            nn.Module(hidden_dim, out_dim),
        )

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.net(x)

batch_size = 256
my_net = SingleLayerNetwork(2, 32, 1).to("cuda")
output = my_net(torch.randn(size=(batch_size, 2)).cuda())
```

- Prefer `net()` over `net.forward()`
- Everything (network and its inputs) on the same device!!!

Torch Best Practices

- When in doubt, **assert** is your friend

```
assert x.shape == (B, N), \
    f"Expected shape ({B}, {N}) but got {x.shape}"
```

- Be extra careful with **.reshape/.view**
 - If you use it, assert before and after
 - Only use it to collapse/expand a single dim
 - In Torch, prefer **.flatten()/.permute()/.unflatten()**
- If you do some complicated operation, test it!
 - Compare to a pure Python implementation

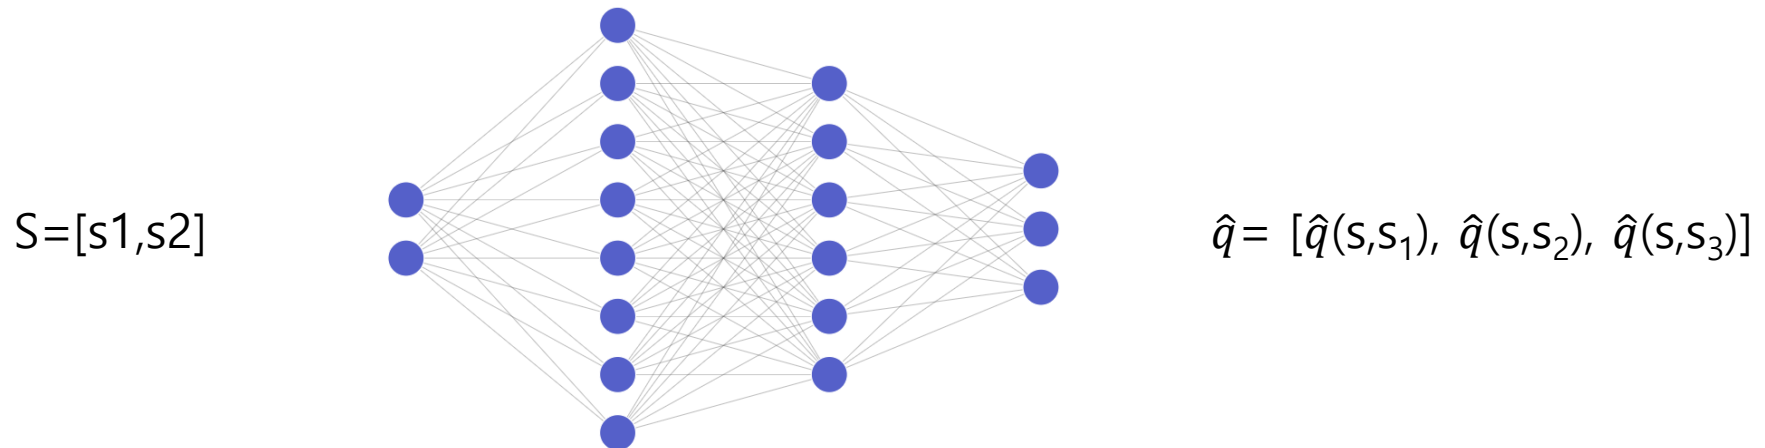
Torch Best Practices

- Don't mix numpy and Torch code
 - Understand the boundaries between the two
 - Make sure to cast 64-bit numpy arrays to 32 bits
 - `torch.Tensor` only in `nn.Module`!
- Training loop will always look the same
 - Load batch, compute loss
 - `.zero_grad()`, `.backward()`, `.step()`

Deep Q-learning

Deep Q-learning

- Q-learning + Neural Network
- Swap the q-value table for a neural network
- Tackle more difficult problems
- Leverage generalization power of neural networks



Neural Net. Architecture for V, Q

- St vector input \rightarrow NN \rightarrow V scalar output
- St, At vector input \rightarrow NN \rightarrow Q scalar output
 - Continuous case
- St vector input \rightarrow Q vector output
 - Discrete action space only
 - Output size is $|A|$

Neural Net. Architecture for policy, π

- S_t input \rightarrow NN \rightarrow vector output
 - Discrete action space case
 - Output size is $|A|$
 - SOFTMAX turns the output into prob. (sum of prob. is 1)
- S_t input \rightarrow NN $\rightarrow \mu_{\theta}(S_t), \delta_{\theta}(S_t)$ output
 - Continuous action space case
 - Represented with Gaussian Distribution

Neural network optimization

- Mean squared error:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N [y_i - \hat{y}_i]^2$$

- We calculate the gradient vector of the cost function with respect to the θ parameters:

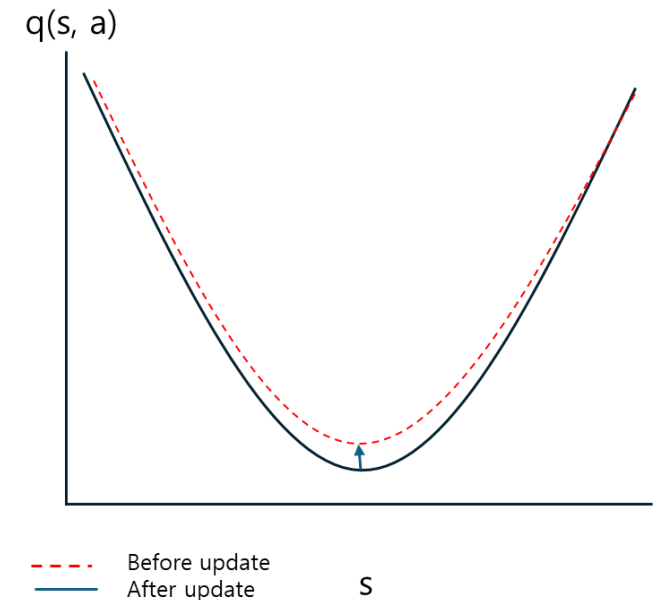
$$\nabla \hat{L}(\theta) = \left[\frac{\partial L}{\partial \theta_1}, \frac{\partial L}{\partial \theta_2}, \dots, \frac{\partial L}{\partial \theta_n} \right]$$

- With the gradient vector, we will make a SGD step:

$$\theta \leftarrow \theta - \alpha \nabla \hat{L}(\theta)$$

Target network

- Two techniques combined generate an unstable learning process.
- Bootstrapping($y_i = R_i + \gamma \hat{q}(S'_i, A'_i | \theta_{targ})$) + function approximator
- Why? : When a value is changed, nearby values will also be affected.
- When a value is changed, nearby values will also be affected.
- By modifying a $\hat{q}(S_i, A_i | \theta)$ estimate we also modify its $\hat{q}(S'_i, A'_i | \theta_{targ})$ target
- For the learning process to be stable, the target must also be stable



Target network

- We make a copy of the neural network to calculate the targets.

$$\theta_{targ} \leftarrow \theta$$

- This neural network does not change with SGD. Its θ parameters remain the same
- The estimated value of S_i', A_i' is calculated with the target network:

$$\hat{L}(\theta) = \frac{1}{N} \sum_{i=0}^N [R_i + \gamma \hat{q}(S_i', A_i' | \theta_{targ}) - \hat{q}(S_i, A_i | \theta)]^2$$

Deep Q-learning

Algorithm 1 Deep Q-Learning

```
1: Input:  $\alpha$  learning rate,  $\epsilon$  random action probability,  
2:    $\gamma$  discount factor,  
3: Initialize q-value parameters  $\theta$  and target parameters  $\theta_{targ} \leftarrow \theta$   
4:  $b \leftarrow \epsilon$ -greedy policy w.r.t  $\hat{q}(s, a|\theta)$   
5:  $\pi \leftarrow$  greedy policy w.r.t  $\hat{q}(s, a|\theta)$   
6: Initialize replay buffer  $B$   
7: for episode  $\in 1..N$  do  
8:   Restart environment and observe the initial state  $S_0$   
9:   for  $t \in 0..T - 1$  do  
10:    Select action  $A_t \sim b(S_t)$   
11:    Execute action  $A_t$  and observe  $S_{t+1}, R_{t+1}$   
12:    Insert transition  $(S_t, A_t, R_{t+1}, S_{t+1})$  into the buffer  $B$   
13:     $K = (S, A, R, S') \sim B$   
14:    Select actions  $A' \sim \pi(S')$   
15:    Compute loss function over the batch of experiences:
```

$$L(\theta) = \frac{1}{|K|} \sum_{i=1}^{|K|} [R_i + \gamma \hat{q}(S'_i, A'_i|\theta_{targ}) - \hat{q}(S_i, A_i|\theta)]^2 \quad (1)$$

```
16:   end for  
17:   Every  $k$  episodes synchronize  $\theta_{targ} \leftarrow \theta$   
18: end for  
19: Output: Near optimal policy  $\pi$  and q-value approximations  $\hat{q}(s, a|\theta)$ 
```

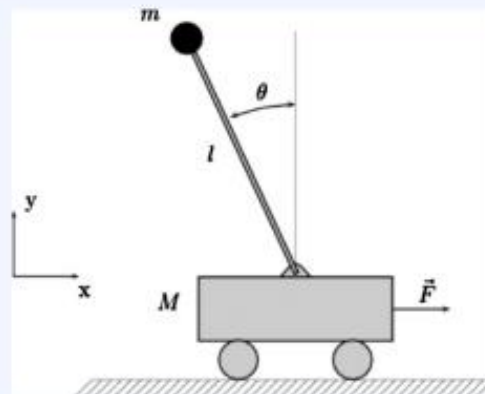
Code Ex.

- Cartpole: Keep the tip of the pole straight

- States: Pole angle and angular velocity
- Actions: Move left, right
- Transition model:

$$\begin{bmatrix} \dot{x} \\ \ddot{x} \\ \dot{\theta} \\ \ddot{\theta} \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \frac{-ml\sin(\theta)\dot{\theta}^2 + F + mg\cos(\theta)\sin(\theta)}{M + m - m\cos(\theta)^2} \\ \dot{\theta} \\ \frac{-ml\cos(\theta))\sin(\theta)\dot{\theta}^2 + F\cos(\theta) + mg\sin(\theta) + Mg\sin(\theta)}{l*(M + m - m\cos(\theta)^2)} \end{bmatrix}$$

- Rewards:
 - +1: Survive
 - 0: Fall



If we know angle and velocity, we can calculate accelerations (Able to describe the whole dynamics)

Code Ex.

The state

The states of the cartpole task will be represented by a vector of four real numbers:

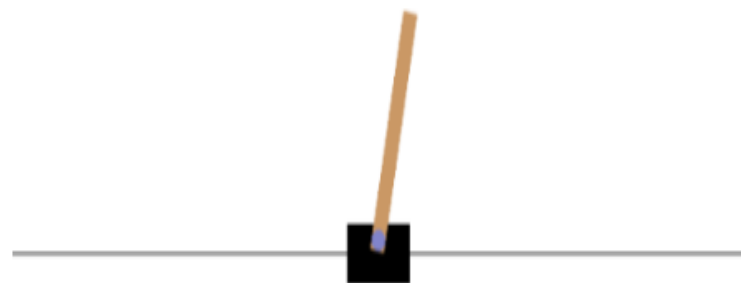
Num	Observation	Min	Max
0	Cart Position	-4.8	4.8
1	Cart Velocity	-Inf	Inf
2	Pole Angle	-0.418 rad (-24 deg)	0.418 rad (24 deg)
3	Pole Angular Velocity	-Inf	Inf

The actions available

We can perform two actions in this environment:

- 0 Apply +1 torque on the joint between the links.
- 1 Do nothing
- 2 Apply -1 torque on the joint between the links.

Reward: +1 for every step



Code Ex.

```
env = gym.make('CartPole-v0')  
#seed_everything(env)  
env.reset()  
plt.imshow(env.render(mode='rgb_array'))
```

```
state_dims = env.observation_space.shape[0]  
num_actions = env.action_space.n  
print(f"CartPole env: State dimensions: {state_dims}, Number of actions: {num_actions}")
```

```
CartPole env: State dimensions: 4, Number of actions: 2
```

Code Ex.

Create the Q-Network: $\hat{q}(s, a|\theta)$

```
q_network = nn.Sequential(  
    nn.Linear(state_dims, 128),  
    nn.ReLU(),  
    nn.Linear(128, 64),  
    nn.ReLU(),  
    nn.Linear(64, num_actions))
```

액션의 개수만큼 output이 나옴

Create the target Q-Network: $\hat{q}(s, a|\theta_{targ})$

```
target_q_network = copy.deepcopy(q_network).eval()
```

Code Ex.

Unsqueeze(dim=0) put a extra. dim. in front.
View(1,-1) put a extra. dim. in front.

For every step,

- **Sample state:**

tensor([[-0.0220, -0.0468, 0.0114, -0.0126]])

- **Next state:**

tensor([[-0.0230, -0.2421, 0.0112, 0.2836]])

- **Reward:**

tensor([[1.]])

- **Done:**

tensor([[False]])

```
class PreprocessEnv(gym.Wrapper):
```

```
    def __init__(self, env):  
        gym.Wrapper.__init__(self, env)
```

```
    def reset(self):  
        obs = self.env.reset()  
        return torch.from_numpy(obs).unsqueeze(dim=0).float()
```

```
    def step(self, action):  
        action = action.item()  
        next_state, reward, done, info = self.env.step(action)  
        next_state = torch.from_numpy(next_state).unsqueeze(dim=0).float()  
        reward = torch.tensor(reward).view(1, -1).float()  
        done = torch.tensor(done).view(1, -1)  
        return next_state, reward, done, info
```

Experience Replay

$[[S_1, A_1, R_2, S_2], [S_2, A_2, R_3, S_3]]$

$[[S_1, S_2], [A_1, A_2], [R_2, R_3], [S_2, S_3]]$

$[[[S_1], [S_2]], [[A_1], [A_2]], [[R_2], [R_3]], [[S_2], [S_3]]]$

A batch shape : N x D for each

Torch.cat은 default가 행(0) 방향으로 붙임
S, A, R, 에 대해 각각 batch(N x D) 로 바꿔 줌

```
class ReplayMemory:

    def __init__(self, capacity=1000000):
        self.capacity = capacity
        self.memory = []
        self.position = 0

    def insert(self, transition):
        if len(self.memory) < self.capacity:
            self.memory.append(None)
        self.memory[self.position] = transition
        self.position = (self.position + 1) % self.capacity

    def sample(self, batch_size):
        assert self.can_sample(batch_size)
        batch = random.sample(self.memory, batch_size)
        batch = zip(*batch)
        return [torch.cat(items) for items in batch]

    def can_sample(self, batch_size):
        return len(self.memory) >= batch_size * 10

    def __len__(self):
        return len(self.memory)
```

Deep Q-learning

$\pi \leftarrow$ greedy policy w.r.t $\hat{q}(s, a|\theta)$

A batch shape : N x D.

$[[[S_1], [S_2]], [[A_1], [A_2]], [[R_2], [R_3]], [[S_2], [S_3]]]$

A batch shape : N x D for each

각 state당 액션의 개수만큼 output이 나옴.

예를들어 $\hat{q}(S_i|\theta)$ 는 $[(\hat{q}(s_i) \text{ for } A_1)], [\hat{q}(s_i) \text{ for } A_2]$

$\hat{q}(S_i, 'A_i'|\theta_{\text{targ}})$, where $A_i' \sim \pi(S')$

$$\hat{L}(\theta) = \frac{1}{N} \sum_{i=0}^N [R_i + \gamma \hat{q}(S_i', A_i'|\theta_{\text{targ}}) - \hat{q}(S_i, A_i|\theta)]^2$$

$$\nabla \hat{L}(\theta) = \left[\frac{\partial L}{\partial \theta_1}, \frac{\partial L}{\partial \theta_2}, \dots, \frac{\partial L}{\partial \theta_n} \right]$$

$$\theta \leftarrow \theta - \alpha \nabla \hat{L}(\theta)$$

```
def deep_q_learning(q_network, policy, episodes,
                    alpha=0.0001, batch_size=32, gamma=0.99, epsilon=0.2):

    optim = AdamW(q_network.parameters(), lr=alpha)
    memory = ReplayMemory()
    stats = {'MSE Loss': [], 'Returns': []}

    for episode in tqdm(range(1, episodes + 1)):
        state = env.reset()
        done = False
        ep_return = 0
        while not done:
            action = policy(state, epsilon)
            next_state, reward, done, _ = env.step(action)

            memory.insert([state, action, reward, done, next_state])

            if memory.can_sample(batch_size):
                state_b, action_b, reward_b, done_b, next_state_b = memory.sample(batch_size)
                qsa_b = q_network(state_b).gather(1, action_b)

                next_qsa_b = target_q_network(next_state_b)
                next_qsa_b = torch.max(next_qsa_b, dim=-1, keepdim=True)[0]

                target_b = reward_b + ~done_b * gamma * next_qsa_b
                loss = F.mse_loss(qsa_b, target_b)
                q_network.zero_grad()
                loss.backward()
                optim.step()

                stats['MSE Loss'].append(loss)

            state = next_state
            ep_return += reward.item()

        stats['Returns'].append(ep_return)

        if episode % 10 == 0:
            target_q_network.load_state_dict(q_network.state_dict())

    return stats
```


Deep Q-learning

- You can cut a gradient by calling `y.detach()`, which will return a new tensor with `requires_grad=False`. Note that `detach` is not an in-place operation! You would want this during evaluation.
- You also can't convert a tensor with `requires_grad=True` to numpy (for the same reason as above). Instead, you need to detach it first, e.g. `y.detach().numpy()`
- `y.clone()` vs. `y.detach()`:
- If `y` change, `y.clone()` not change.
- If `y` change `y.detach()` change

Create the exploratory policy: $b(s)$

```
def policy(state, epsilon=0.):  
    if torch.rand(1) < epsilon:  
        return torch.randint(num_actions, (1, 1))  
    else:  
        av = q_network(state).detach()  
        return torch.argmax(av, dim=-1, keepdim=True)
```

Code Ex.

- Open `deep_q_learning_colab.ipynb`