

# 11강. Policy Gradient methods

# Contents

- Policy Gradient
  - Overview
  - Advantages
  - Stochastic policy
  - Policy performance
  - Stochastic Gradient Ascent
  - Policy Gradient Theorem
- REINFORCE
  - Overview
  - Parallel learning
  - Exploration strategy
  - Code Ex. REINFORCE + Parallel Agent Learning
- Actor-Critic
  - Overview
- A2C
  - Overview
  - Return vs. Advantage
  - Stochastic Gradient Ascent
  - REINFORCE vs. A2C
  - Code Ex.: A2C + Parallel Agent Learning


# Policy Gradient methods

# Overview

- So far, we learned Q value-based methods, the policy is defined based on Q values


Value table:

$$a = \arg \max_a Q(s, a)$$



Function approximators:

$$a = \arg \max_a \hat{q}(s, a | \theta)$$




# Overview

- PG methods use a function approximator, not to estimate Q value, but to estimate the probabilities of taking each action:

$$\pi(a|s, \theta) \in [0,1]$$

- The neural network is the policy (stochastic):

Input:  
 $s=[s_1, s_2]$



Output:  
 $\pi(S|\theta) = [p(a_1), p(a_2), p(a_3)]$

# Advantages

- Value-based methods cannot represent stochastic policies in a simple way

Greedy policy:

$$\pi(a'|s) \begin{cases} 1 & \text{if } a' = \arg \max_a \hat{q}(s, a|\theta) \\ 0 & \text{else} \end{cases}$$

Epsilon-greedy policy

$$(a'|s) \begin{cases} 1 - \varepsilon + \frac{\varepsilon}{|A|} & \text{if } a' = \arg \max_a \hat{q}(s, a|\theta) \\ \frac{\varepsilon}{|A|} & \text{else} \end{cases}$$

- Imagine a poker game where the agent has imperfect information  
If  $\pi^*(s) = [0.7, 0.3]$ , how do we represent it?

# Advantages

- The policy changes more smoothly during learning:

## Value-based methods:

when the maximum q-value changes, they choose a new action 100% of the time.

$$a = \arg \max_a \hat{q}(s, a | \theta)$$

## Policy gradient methods:

the probability of choosing an action changes in small increments

$$a \sim \pi(s | \theta)$$


The probabilities of taking an action gradually increase if the action turns out to be effective and gradually decrease if not

# Stochastic policy

- The neural network can be viewed as a function:

$$y = \varphi_2(\varphi_1(x \cdot w1) \cdot w2)$$

$$W1 = \begin{bmatrix} w11 & w12 & w13 & w14 & w15 & w16 \\ w21 & w22 & w23 & w24 & w25 & w26 \\ w31 & w32 & w33 & w34 & w35 & w36 \end{bmatrix} \quad W2 = \begin{bmatrix} w11 & w12 \\ w21 & w22 \\ w31 & w32 \\ w41 & w42 \\ w51 & w52 \\ w61 & w62 \end{bmatrix}$$




```
nural_network = nn.Sequential(  
    nn.Linear(3, 6),  
    nn.ReLU(),  
    nn.Linear(6, 2))
```

- By changing its parameters  $W1$ , we can modify it to approximate the function we are interested in



# Stochastic policy

- By normalizing the neural network's output, **softmax activation function** is perfect tool to generate probability vector,  $y$



$$y = [\sigma(x_1), \sigma(x_2), \dots, \sigma(x_3)]$$

e. g.,  $y = [0.6, 0.39, 0.01]$

$$\sigma(x_i) = \frac{e^{x_i}}{\sum e^{x_i}}$$

$$\sum_i \sigma(x_i) = 1$$

# Stochastic policy

- Discrete action space:



- Continuous action space:



# Policy performance

- If we want to find the optimal policy, we need to be able to compare them.
- We will define policy performance as:  $J(\theta)$  the performance of the policy is a function of its parameters

$$\begin{aligned} J(\theta) &= E^{\pi^\theta}[R|s] \\ &= V^{\pi^\theta}(s) \\ &= \sum_a \pi_\theta(a|s) Q^{\pi^\theta}(s, a) \end{aligned}$$

By changing the parameters of the neural net., we change the policy

# Policy performance

- If  $J_{\pi_1}(\theta) > J_{\pi_2}(\theta)$ , then we consider that  $J_{\pi_1}(\theta)$  is better than  $J_{\pi_2}(\theta)$

- Our goal is to find the  $\theta$  values that maximize policy performance:

$$\pi * (a|s, \theta) = \arg \max_a J(\theta)$$

- We will use the experience samples that the agent obtains to approximate the policy's performance as  $\hat{J}(\theta)$

# SGA

- We will approximate the optimal  $\theta$  values by **stochastic gradient ascent (SGA)**:

$$\theta_{t+1} = \theta_t + \underline{\alpha \nabla \hat{J}(\theta)}$$

we'll take the gradient step  
in the direction of steepest  
ascent

where:

$$\nabla \hat{J}(\theta) = \left[ \frac{\partial \hat{J}(\theta)}{\partial \theta_1}, \frac{\partial \hat{J}(\theta)}{\partial \theta_2}, \dots, \frac{\partial \hat{J}(\theta)}{\partial \theta_n} \right]$$

# SGA

- Optimizing the policy using values that the agent can observe.

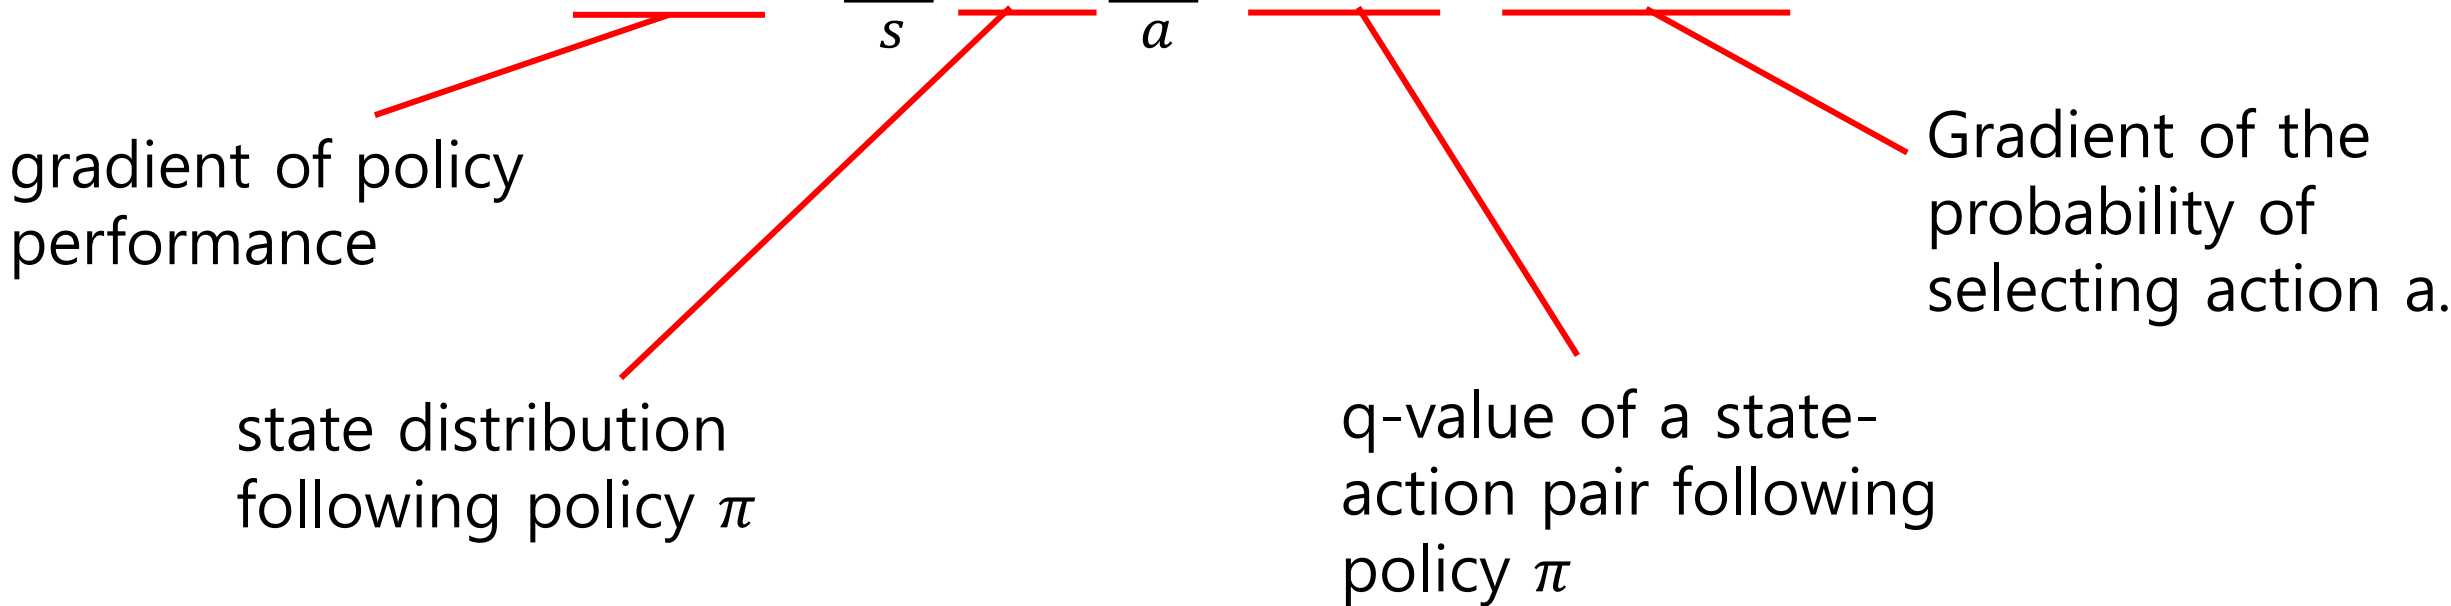
$$\begin{aligned} \bullet \nabla_{\theta} J(\theta) &= \sum_a \nabla_{\theta} \pi_{\theta}(a|s) Q^{\pi_{\theta}}(s, a) \\ &= \sum_a \underbrace{\pi_{\theta}(a|s)}_{\text{Distribution of collected data}} Q^{\pi_{\theta}}(s, a) \underbrace{\nabla_{\theta} \log \pi_{\theta}(a|s)}_{\text{Equal to } \frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi_{\theta}(a|s)}} \end{aligned}$$

$$= \mathbb{E}_a^{\pi_{\theta}} [Q^{\pi_{\theta}}(s, a) \nabla_{\theta} \log \pi_{\theta}(a|s)]$$

$$\approx \underline{\hat{E}} [Q^{\pi_{\theta}}(s, a) \nabla_{\theta} \log \pi_{\theta}(a|s)]$$

Sample  
mean that  
follows  $\pi_{\theta}$

# Policy gradient theorem

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a|s, \theta)$$


The diagram illustrates the components of the policy gradient theorem equation. Red lines connect the terms as follows: a line from  $\nabla J(\theta)$  to 'gradient of policy performance'; a line from  $\mu(s)$  to 'state distribution following policy  $\pi$ '; a line from  $q_\pi(s, a)$  to 'q-value of a state-action pair following policy  $\pi$ '; and a line from  $\nabla \pi(a|s, \theta)$  to 'Gradient of the probability of selecting action a.'

gradient of policy performance

state distribution following policy  $\pi$

q-value of a state-action pair following policy  $\pi$

Gradient of the probability of selecting action a.

# Policy gradient theorem

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a|s, \theta)$$

- The PG is,
  - proportional to the return of each action in each state,
  - multiplied by the gradient of the probability of taking that action in that state
  - and weighted by the frequency with which we observe each state following that policy



# Policy gradient theorem

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a|s, \theta)$$

- If an action in a state produces a **positive return**, we must **increase the probability** of picking that action to increase the return.
- If that action produces a **negative return**, the probability of taking that action must be **reduced**.

REINFORCE

# Overview

- Policy gradient + Monte Carlo
- We will perform stochastic gradient ascent(SGA):

$$\theta_{t+1} = \theta_t + \alpha \nabla \hat{J}(\theta)$$

- To do this we will have to approximate the gradient of the policy performance estimate  $\nabla \hat{J}(\theta)$  using samples collected from the environment

# Overview

- Note there's no bootstrapping,  $Q^{\pi\theta}(s, a) = E^{\pi\theta}[R|s, a]$

- $$\begin{aligned}\nabla_{\theta} J(\theta) &= \sum_a \nabla_{\theta} \pi_{\theta}(a|s) \overbrace{R_{s,a}}^{\text{Distribution of collected data}} \\ &= \sum_a \underbrace{\pi_{\theta}(a|s)}_{\text{Distribution of collected data}} \overbrace{R_{s,a} \nabla_{\theta} \log \pi_{\theta}(a|s)}^{\text{Equal to } \frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi_{\theta}(a|s)}}\end{aligned}$$

$$= E_a^{\pi\theta} [\overbrace{R_{s,a}}^{\text{Distribution of collected data}} \nabla_{\theta} \log \pi_{\theta}(a|s)]$$

$$\approx \underline{\hat{E}} [\overbrace{R_{s,a}}^{\text{Distribution of collected data}} \nabla_{\theta} \log \pi_{\theta}(a|s)]$$

Sample  
mean that  
follows  $\pi\theta$

# Overview

- We will perform stochastic gradient ascent(SGA):

$$\theta_{t+1} = \theta_t + \alpha \nabla \hat{J}(\theta)$$

- REINFORCE perform stochastic gradient ascent(SGA):

$$\theta_{t+1} = \theta_t + \alpha \gamma^t G_t \nabla \ln \pi(a|s, \theta)$$

# Overview

Initialize  $\theta$  of the parameterized policy  $\pi_\theta$  and learning rate  $\eta$

While True

Generate an episode  $e = (s_0, a_0, r_1, s_1, \dots, s_{T-1}, a_{T-1}, r_T, s_T)$  using  $\pi_\theta$

- Get an episode

$$\Delta_\theta = 0$$

For  $t \in 0 : T - 1$

$$g_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_{k+1}$$

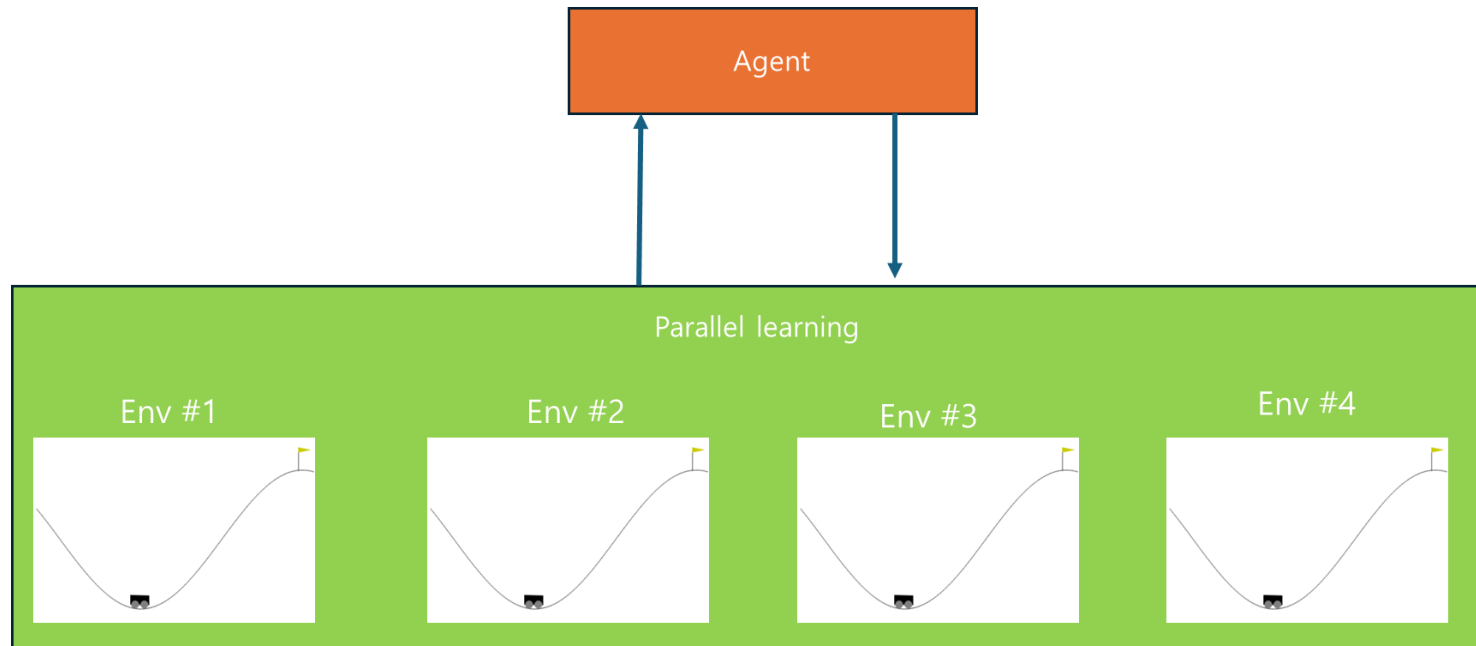
$$\Delta_\theta \leftarrow \Delta_\theta + g_t \nabla_\theta \log \pi_\theta(a|s)$$

$$\theta \leftarrow \theta + \eta \cdot \Delta_\theta$$


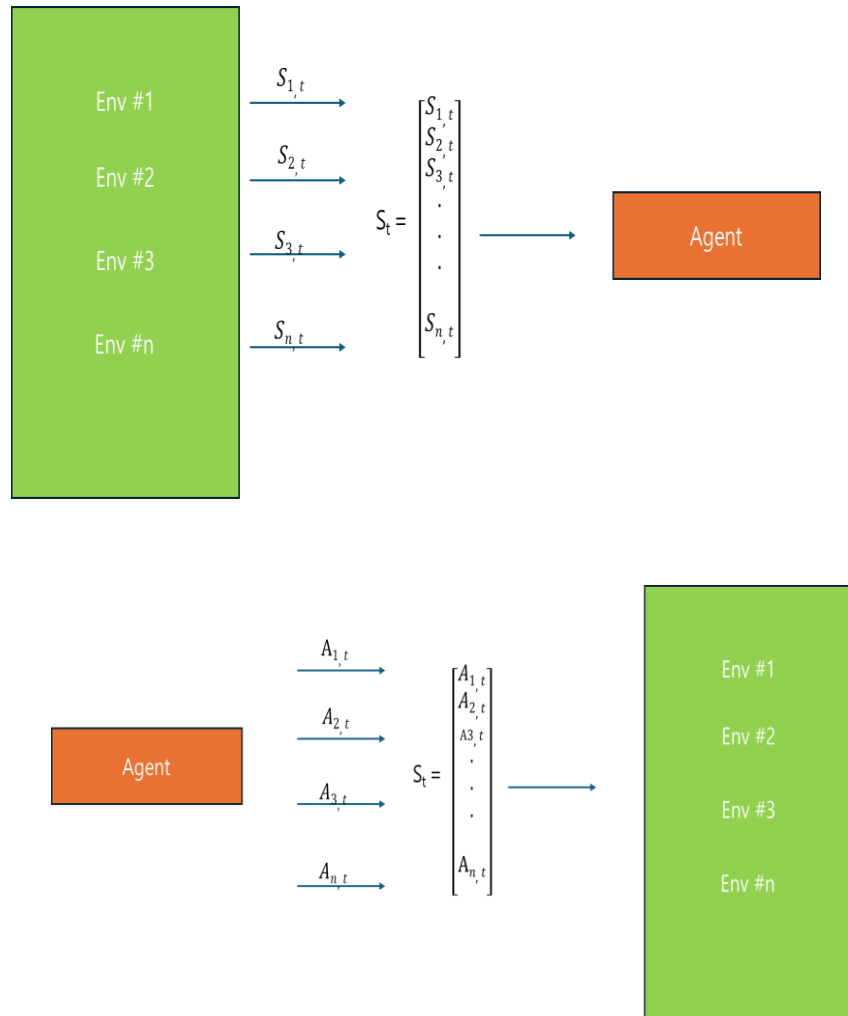
- $g_t$  is the actual return obtained starting at time  $t$  of the episode (all the rewards obtained until the end of the episode)
- the gradient of the policy performance estimate  $\Delta_\theta$  using samples collected from the environment
- Stochastic gradient ascent:

# Parallel learning

- Successive state tend to be very similar to the previous one, and this is called the time correlation problem
- The solution used with Policy Gradient methods



# Parallel learning





# Exploration Strategy

- We want to maintain the agent's exploration but we do not have mechanisms such as  $\varepsilon$ -greedy policies
- Now the neural net. is our policy. How to incorporate an exploration mechanism into our neural network?
- We will incentivize the agent to keep the entropy of its policy as high as possible

$$H(X) = - \sum_{x \in X} p(x) \cdot \ln p(x)$$


# Exploration Strategy

- What is Entropy? the level of uncertainty of a random variable

e.g.,

if  $p(X = x_1) = 1, p(X = x_2) = 0, H(X) = -[1 \cdot \ln(1) + 0 \cdot \ln(0)] = 0$

if  $p(X = x_1) = 0.5, p(X = x_2) = 0.5, H(X) = -[0.5 \cdot \ln(0.5) + 0.5 \cdot \ln(0.5)] \approx 0.6931$



- Max. point where our confidence in our predictions will go down.
- Max. point where it surprises us

# Exploration Strategy


- What is the entropy of a policy?
- Uncertainty in the action to be selected in a state:

$$H_{\pi}(A_t) = - \sum_{a \in A_t} \pi(a|St) \cdot \ln \pi(a|St)$$

- Random variable( $A_t$ ) is the action the policy will choose in a state.
- The entropy is computed by multiplying the probability of choosing each action by its logarithm and adding up the results

# Exploration Strategy

- What is the entropy of a policy?
- Imagine that we have 4 actions available



# Exploration Strategy

- What is the entropy of a policy?
- In order for the agent to explore the environment. We must keep the entropy of the policy high

- We add the entropy to the function to be maximized:

$$\theta_{t+1} = \theta_t + \alpha [ \gamma^t G_t \nabla \ln \pi(a|s, \theta) + \beta \nabla H(\pi) ]$$

- Advantages in optimizing a policy:  
Exploration, Robustness, Policy refinement

# Parallel REINFORCE learning

---

**Algorithm 1** REINFORCE

---


```
1: Input:  $\alpha$  learning rate,  $\gamma$  discount factor.
2: Initialize parallel environments  $E$ 
3: Initialize policy parameters  $\theta$ 
4: for episode in 1..N do
5:   Use  $\pi(s|\theta)$  to collect  $|E|$  trajectories:  $S_0, A_0, R_1, \dots, R_T$ 
6:    $G = \vec{0}$ 
7:   for  $t = T-1..0$  do
8:      $G = R_t + \gamma G$ 
9:     Compute entropy regularization:  $H_t = -\sum_a \pi(a|S_t) \ln \pi(a|S_t)$ 
10:     $\hat{J}(\theta) = \gamma^t G \ln \pi(A_t|S_t, \theta) + H_t$ 
11:     $\theta = \theta + \alpha \nabla \hat{J}(\theta)$ 
12:   end for
13: end for
```

---

We also try to maximize the entropy of the policy


# Code Ex.

- Upload 'REINFORCE\_CartPole.ipynb' file onto Colab
- Upload 'utils.py' file onto Colab
- Upload 'parallel\_env.py' file onto Colab
- Add 'pip install numpy==1.23.1'



# Code Ex.

- Cartpole: move a cart (black) such that it balances a pendulum (brown) without moving too far from the center.
- State: The agent observes current position and velocity of the cart, as well as angle and velocity of the pole (cart position, cart velocity, pole angle, pole angular velocity)
- Action: It can act by pushing the cart to the left (value 0) or to the right (value 1).
- Reward: +1 for every step







# Actor-Critic

# Overview

- 'Monte-Carlo' Policy Gradient method
  - $J(\theta) = \sum_a \pi_\theta(a|s) R_{s,a}$
- 'Actor-Critic' method requires Q or V to calculate PG
  - $J(\theta) = \sum_a \pi_\theta(a|s) Q^{\pi_\theta}(s, a)$
  - $J(\theta) = \sum_{s'} \sum_a p(s'|s, a) \pi_\theta(a|s) [r + \gamma V^{\pi_\theta}(s)]$

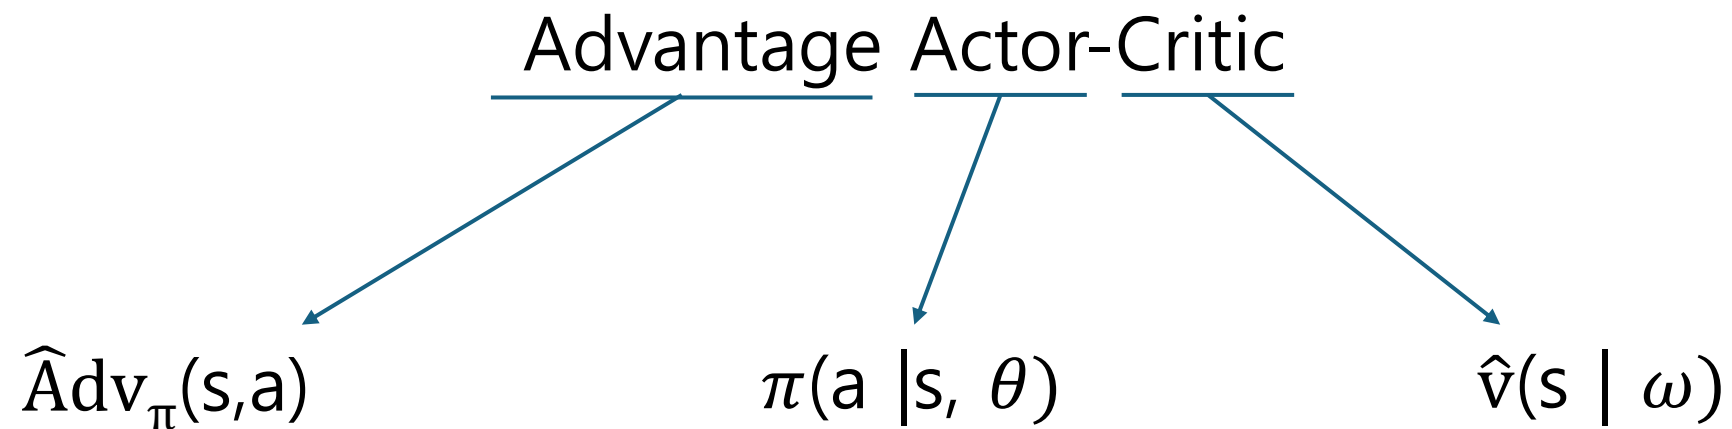
# Overview

- TD + Policy Gradient
  - Optimize the neural network during the episode by bootstrapping the value of the next state
  - Recall Temporal Difference (TD) learn at each time step
- Agent has both actor and critic networks



Advantage Actor-Critic(A2C)

# Overview



# Overview

- Excess return of choosing the  $a$  action instead of following the policy:

$$\text{Adv}_{\pi}(s,a) = q_{\pi}(s, a) - v_{\pi}(s)$$

$$\text{Adv}_{\pi}(s,a) = q_{\pi}(s, a) - \sum_a \pi(a|s)q_{\pi}(s, a)$$

- We want to reinforce the actions that obtain better results and discourage those that obtain worse results

# Overview

$$\text{Adv}_{\pi}(s,a) = q_{\pi}(s, a) - v_{\pi}(s)$$

$$\text{Adv}_{\pi}(s,a) > 0 :$$

Taking the action,  $a$  is **better** than simply following the policy

$$\text{Adv}_{\pi}(s,a) < 0 :$$

Taking the action  $a$  is **worse** than following the policy


- We can estimate the advantage of an action as:

$$\hat{\text{Adv}}_{\pi}(s,a) = r(s, a) + \gamma \hat{v}_{\pi}(s') - \hat{v}_{\pi}(s)$$

# Update rule: REINFORCE

Assume agent gets -1  
reward each time spent

- REINFORCE:  $\theta_{t+1} = \theta_t + \alpha \gamma^t \underline{G_t} \nabla \ln \pi(A_t | S_t, \theta_t)$   
The best is -1



- If  $G_t < 0$ , discourages the actions that will have negative return
  - If an action led to a negative return, that action was disincentivized **even if it was the least bad action available**
  - Eventually it gets to the best policy, but a long time to get there
  - Using the return, the agent was not able to make these distinctions




# Update rule: A2C

- A2C:

$$\theta_{t+1} = \theta_t + \alpha \gamma^t \underbrace{\widehat{\text{Adv}}_t}_{\text{Adv}(s, \downarrow)} \nabla \ln \pi(A_t | S_t, \theta_t)$$

- If  $\widehat{\text{Adv}}_t > 0$ , reinforces the actions that are better
  - In a bad state, the advantage function helps the agent understand which action is the lesser evil and helps it to reinforce that action
  - If action,  $\downarrow$  is lesser evil,  $\text{Adv}(s, \downarrow)$  is positive
  - It also allows it to see which action is the worst in a very good state and disincentivize it



# SGA

- Where the advantage is the temporal difference error

$$\hat{\text{Adv}}_t = R_{t+1} + \gamma \hat{v}(s_{t+1} | \omega) - \hat{v}(s_t | \omega)$$

- Bootstrapped value. We don't need to the end of episode to update the NN
- Bootstrapped value: the estimate of the rest of the rewards that we expect to get starting from the next state
- The baseline eliminates the effect of the state that we're in.
- Subtract this because we want to evaluate only the merit of that action

# REINFORCE vs. A2C

- REINFORCE:
- Used true return:  $R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t-1} R_T$
- Unbiased empirical return, but with high variance
- A2C:
- Used bootstrapping:  $R_{t+1} + \gamma \hat{v}(s_{t+1}) - \hat{v}(s_t)$ 
  - Learning occurs during the episode
  - we don't need to the end of episode to update the NN.
- Reducing variance accelerates learning
  - Bootstrapping introduces bias but reduces variance

---

**Algorithm 1** Advantage Actor-Critic (A2C)

---

- 1: **Input:**  $\alpha$  learning rate,  $\gamma$  discount factor.
- 2: Initialize parallel environments  $E$
- 3: Initialize policy network parameters  $\theta$
- 4: Initialize value network parameters  $w$
- 5: **for** episode in 1..N **do**
- 6:     Initialize parallel environments  $E$  and obtain initial states  $S_0$
- 7:     **for**  $t=0..T-1$  **do**
- 8:          $A_t \sim \pi(S_t|\theta)$
- 9:         Execute the actions in the environments  $E$  and obtain  $R_{t+1}, S_{t+1}$
- 10:        Update the value network with SGD:

$$L(w) = \frac{1}{|E|} [R_{t+1} + \gamma v(S_{t+1}|w) - v(S_t|w)]^2 \quad (1)$$

$$\theta = \theta - \alpha \nabla L(w) \quad (2)$$

- 11:     Update the policy network with SGA:

$$Adv_t = R_{t+1} + \gamma v(S_{t+1}|w) - v(S_t|w) \quad (3)$$

$$\hat{J}(\theta) = \gamma^t Adv_t \ln \pi(A_t|S_t, \theta) + H_t \quad (4)$$


$$\theta = \theta + \alpha \nabla \hat{J}(\theta) \quad (5)$$

- 12:     **end for**
  - 13: **end for**
- 

$$H(X) = - \sum_{x \in X} p(x) \cdot \ln p(x)$$

# Code Ex.

- Arcrobot: swing the lower part of a two-link robot up to a given height.
- Observations space: The agent observes current positions and velocities of the joints.
- Action space: It can act by applying positive torque (value 0), no torque (value 1), or negative torque (value 2) only to the joint between the two links.



| 0                | 1                | 2                | 3                | 4                | 5                |
|------------------|------------------|------------------|------------------|------------------|------------------|
| $\cos(\theta_1)$ | $\sin(\theta_1)$ | $\cos(\theta_2)$ | $\sin(\theta_2)$ | $\dot{\theta}_1$ | $\dot{\theta}_2$ |

| Action | 0  | 1 | 2 |
|--------|----|---|---|
| Torque | -1 | 0 | 1 |

# Code Ex.

- Upload 'A2C\_Acrobot.ipynb' file onto Colab
- Upload 'utils.py' file onto Colab
- Upload 'parallel\_env.py' file onto Colab
- Add 'pip install numpy==1.23.1'

